

GAME ENGINE USING C++11

JASON JURECKA

HELLO & WELCOME

Disclaimer: I work at Blizzard but not representing a Blizzard product.

ABOUT ME...

- 12+YEAR VETERAN
- 10+ TITLES/PROJECTS CONTRIBUTED TO.
- HAVE THE RESPONSIBILITIES OF A SWISS ARMY KNIFE.



TABLE OF CONTENTS SUMMARY

- PROJECT MENTALITY
- GAME ENGINE OVERVIEW
- HOW DOES C++11 HELP US
- C++11 STANDARD HIGHLIGHTS
- LANGUAGE NICE TO HAVE THINGS
- QUESTIONS
- BONUS MATERIAL (TIME PERMITTING)

PROJECT MENTALITY

- Mass Concurrency
- Leverage std:: over roll my own
- Keep things simple
- DO NOT get sucked into 'C' style patterns
- DO NOT focus on latest graphics
- Work for 60fps - frame ~ 16ms
- Learn/Grow



keep it simple.

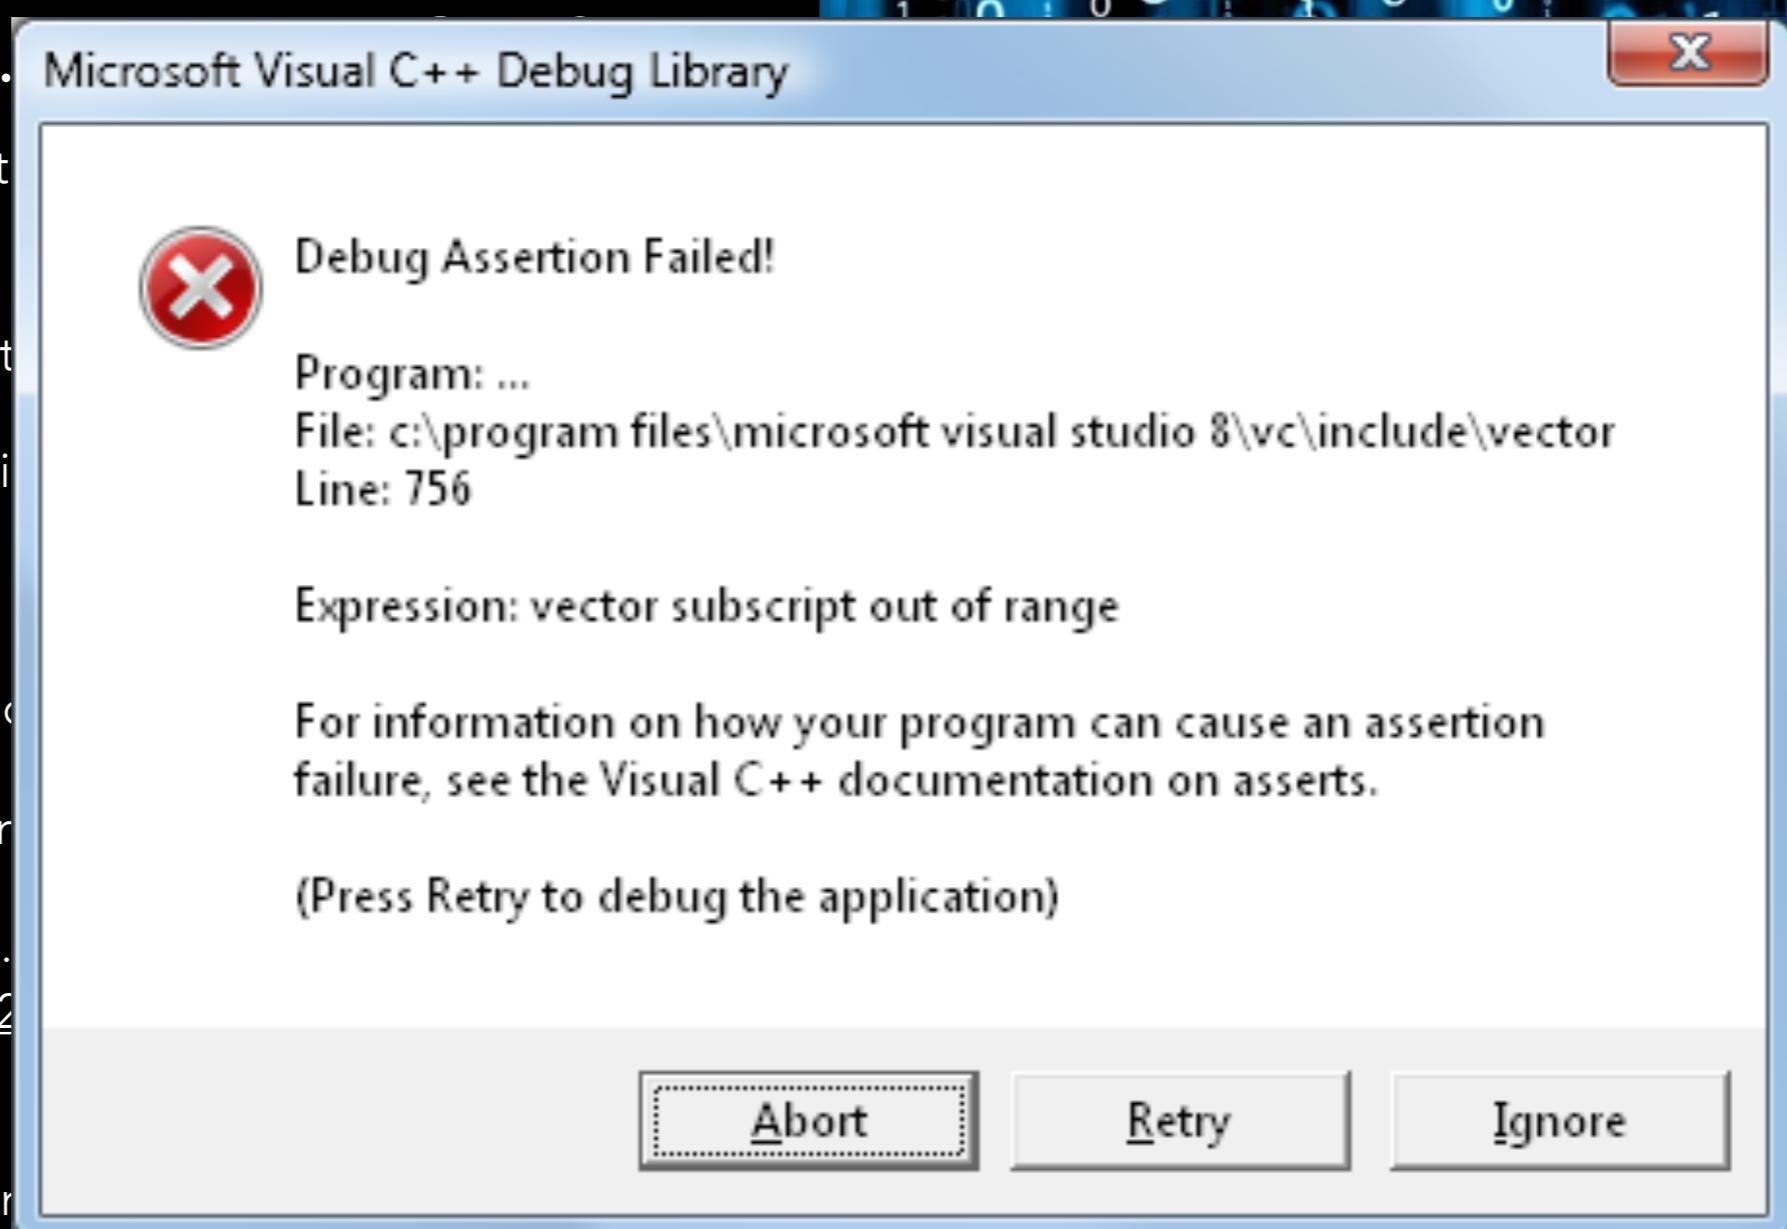
WHAT IS THE MAIN JOB OF A GAME ENGINE?

- Interface with Hardware API's (graphics, sound, network, input...)
- Data creation/loading
- Data interpretation
- Data Presentation



GAME ENGINES DON'T USE STD MUCH...

- C++98 Replacement
years why change").
- Hardware implement
- Containers + std::strin
 - Slow Iteration
 - Slow manipulation
 - Memory usage r
 - Flat Containers .
std.org/jtc1/sc22/p0038r0.html
- Code size add for ter
- <https://www.youtube.com/watch?v=p6NxZ5Qnzfo>



Why Not? (Cont)

- Throwing of exceptions
 - Many game code bases turn them off all together
 - Limit's 3rd Party library usage
 - Some Libraries use exceptions as execution flow
 - QA is supposed to find issues during testing phase
- We can patch to fix found things in the wild
 - It is a game not a cardiac bypass machine/moon rover
- Cost Explained link <http://h-deb.clg.qc.ca/Sujets/Developpement/Exceptions-Costs.html>

WHY NOT? (CONT)

- Not all games are created equal
 - Different concerns if RTS vs FPS vs RPG
 - Same with open world/level based/side-scroller
 - Platforms restrictions (XBOX/PS vs PC/MAC)
 - Console dedicated threads or smaller thread pool
 - PC variant hardware configurations
 - Mobile devices (android/ios)
 - Web browser based limits (HTML5, flash)

HOW DOES C++11 HELP US? - DATA CREATION/LOADING

- Compile time resource versioning using `constexpr`
- Compile time type validation using `static_assert`
- Compile time serialize/deserialize function lists

COMPILE TIME RESOURCE VERSIONING USING CONSTEXPR

- What would go into computing a version hash for this structure?
 - name of structure
 - data member type
 - data member name
 - data member order
 - data alignment

```
struct ·Vector3 ·
{
    → float ·x ·= ·0.0f;
    → float ·y ·= ·0.0f;
    → float ·z ·= ·0.0f;
};

struct ·SpawnPoint ·
{
    → Vector3 ·m _Position;
    → Vector3 ·m _Direction;
};
```

COMPILE TIME RESOURCE VERSIONING USING CONSTEXPR (CONT)

```
using::HashValue::=unsigned::int;

class::constexprString{
→    const::char*·p;
→    std::size_t·sz;
public:
→    template<std::size_t·N>
→    constexpr::constexprString(const::char(&a)[N])·:·p(a),·sz(N-·1){}
→    constexpr::char·operator[](std::size_t·n)·const·{·return·(n·>=·sz)·?·'\0'·:·p[n];·}
→    constexpr::std::size_t·size()·const·{·return·sz;·}
};

//Constant·Expression·using·recursion·to·create·a·hash·value
constexpr::HashValue·GetHash_constexprString(const::constexprString&·_toHash,·const·unsigned·index·=·0)
{
→    return·(·index·>=·_toHash.size())·?·0·:·(static_cast<unsigned>(_toHash[index])·*·(index+1)·\
→    ·*·s_PrimeTable[(static_cast<unsigned>(_toHash[index])·%·s_PrimeNumCount)])·\ \
→    ·+·GetHash_constexprString(_toHash,·index·+·1);
}

template<std::size_t·N>
constexpr::HashValue·GetHash_constexprString_Array(const::constexprString(&_toHash)[N],·const·unsigned·index·=·0)
{
→    return·(index·>=·N)·?·0·:·GetHash_constexprString(_toHash[index],·0)·+·\ \
→    ·(GetHash_constexprString_Array(_toHash,·index·+·1)·<<·index);
}
```

COMPILE TIME RESOURCE VERSIONING USING CONSTEXPR (CONT)

```
struct·SpawnPoint           struct·SpawnPoint           struct·spawnPoint          struct·SpawnPoint
{                           {                           {                           {
    → Vector3·m_Position;   → Vector3·m_Position;   → Vector3·m_Position;   → Vector3·m_Direction;
    → Vector3·m_Direction;   → Vector3·m_direction;  → Vector3·m_Direction;   → Vector3·m_Position;
}};                         };                         };                         };
```

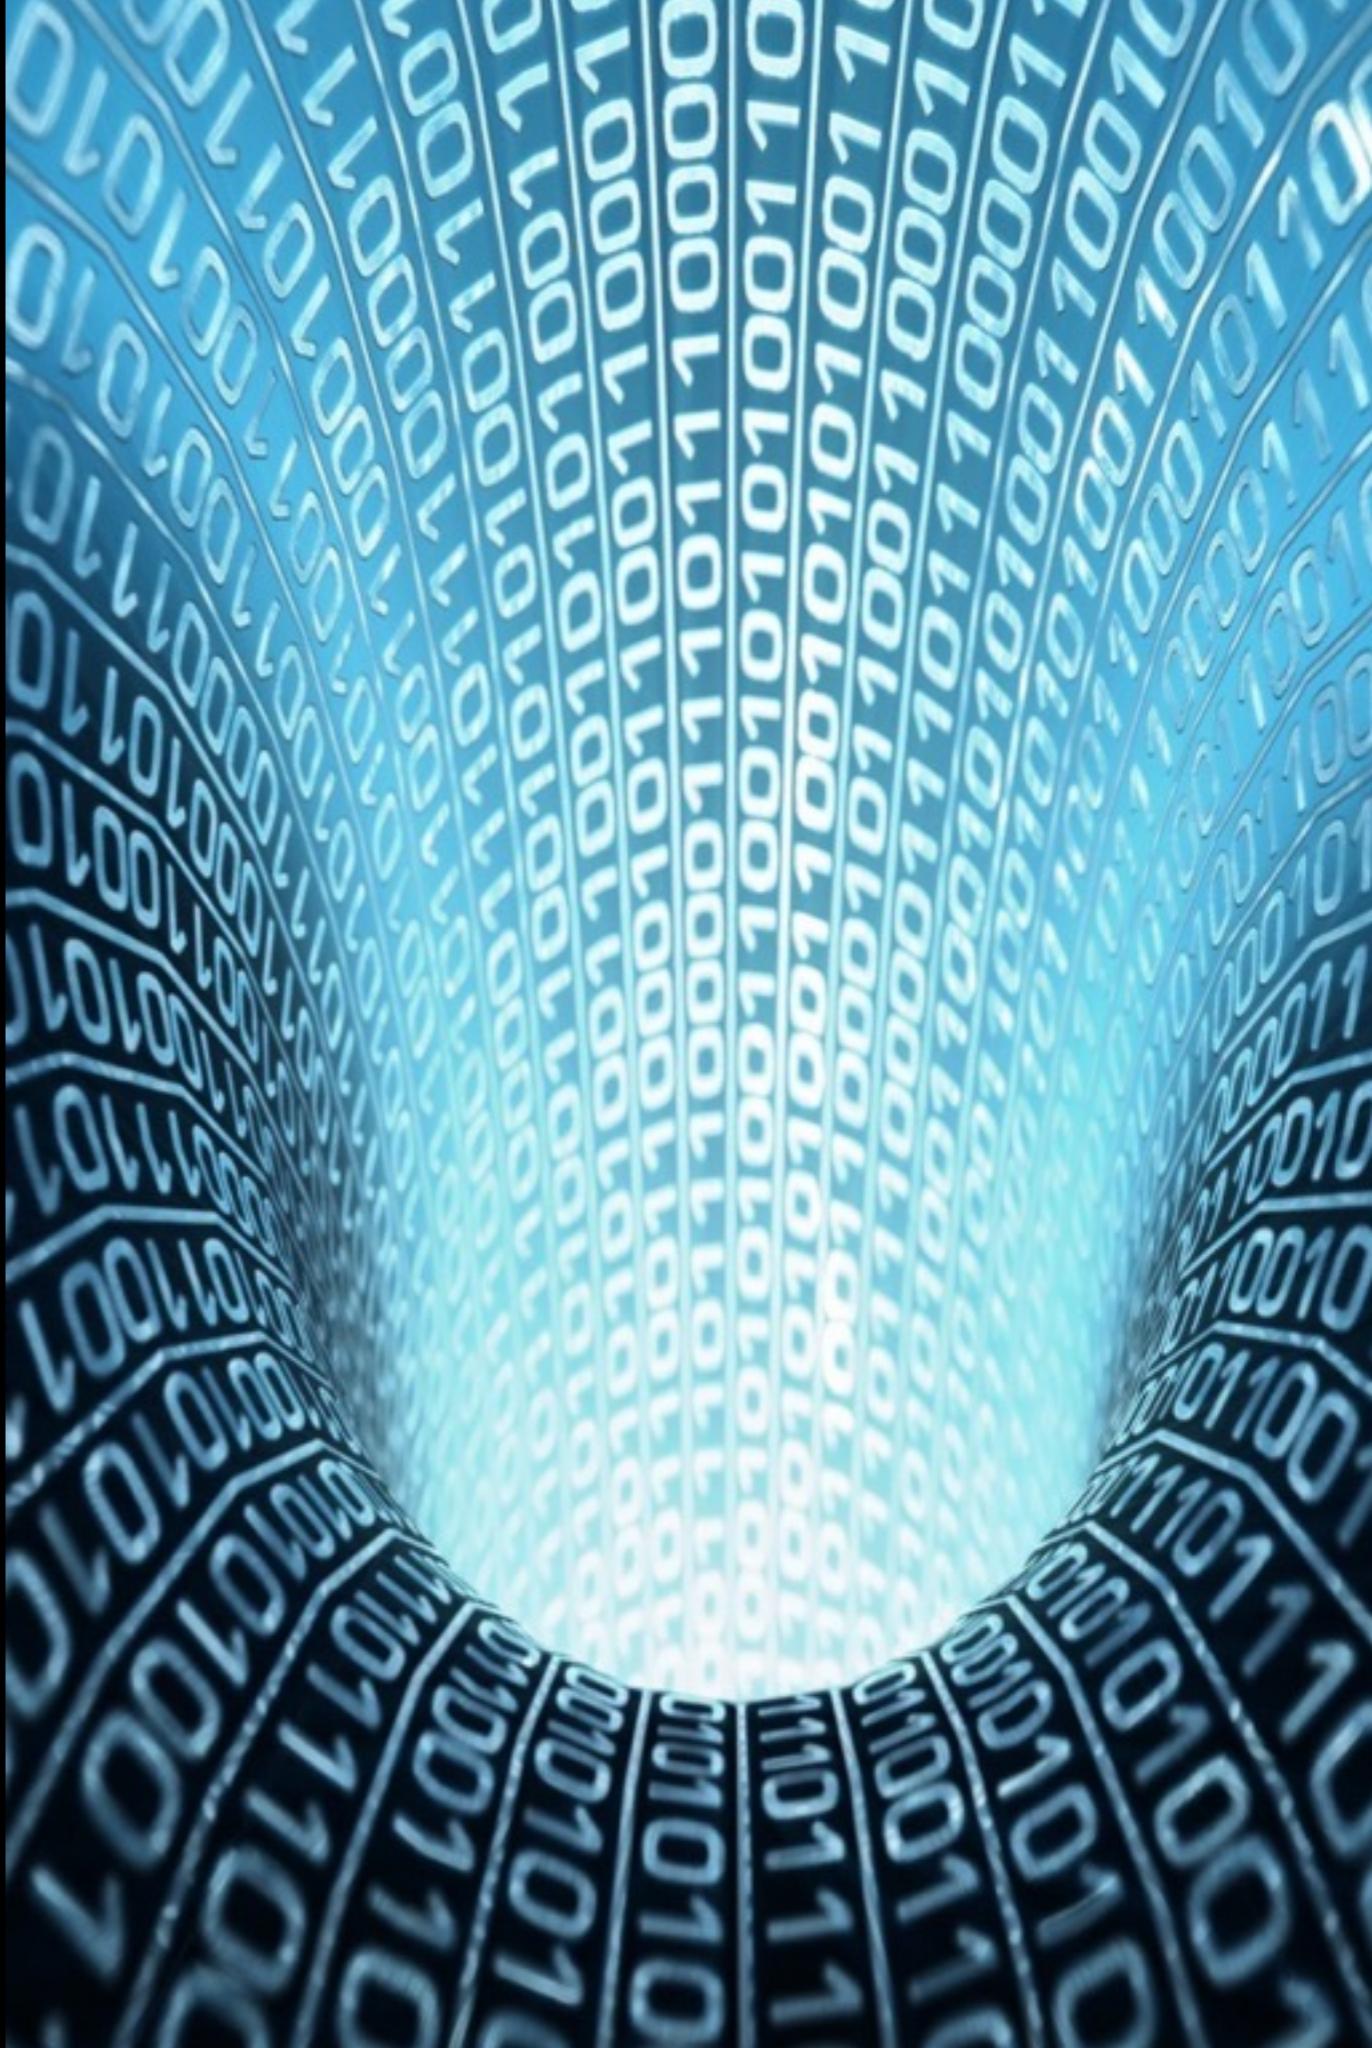
COMPILE TIME RESOURCE VERSIONING USING CONSTEXPR (CONT)

```
RSC_DATA_BEGIN(Vector3)
→ RSC_DATA(Vector3, float, m_X)
→ RSC_DATA(Vector3, float, m_Y)
→ RSC_DATA(Vector3, float, m_Z)
RSC_DATA_END(Vector3)

RSC_DATA_BEGIN(SpawnPoint)
→ RSC_DATA(SpawnPoint, rscVector3, m_Position)
→ RSC_DATA(SpawnPoint, rscVector3, m_Direction)
RSC_DATA_END(SpawnPoint)
```

MASS CONCURRENCY

- Most (if not almost all) devices now have multiple cores and thread support
 - Why not use that power?
- Benefits
 - Exploit multi processor machines.
 - Exploit cloud processing (dynamic scaling)
 - Tasks can be small and simple (focused optimizations)
 - serialized sequences can be distributed
 - highly modular for version-ing of tasks/systems
 - efficiency of doing multiple things at once out weighs serialized process cost.



MASS CONCURRENCY (CONT)

- Concerns:
 - Tasks need to be as independent as possible
 - Data manipulation needs to be thread safe (dead-locks or live-locks)
 - Need to keep in mind when data synchronization happens
 - Need to account for things executing in non determined order (ish)
 - Task queue speed (lock vs lock-free)
 - Thread context switching

TASKSYSTEM

```
namespace TaskMaster
{
    → bool Initialize();
    → void Shutdown();
    → void Process();

    → void AddTask(std::unique_ptr<ITask> task);
    → std::future<bool> AddTrackedTask(std::unique_ptr<ITask> task, ITask* parent);
}
```

TASKSYSTEM (CONT)

```
void TaskMaster::AddTask(std::unique_ptr<ITask> task)
{
    assert(task);
    s_Tasks.push(std::move(task));
}

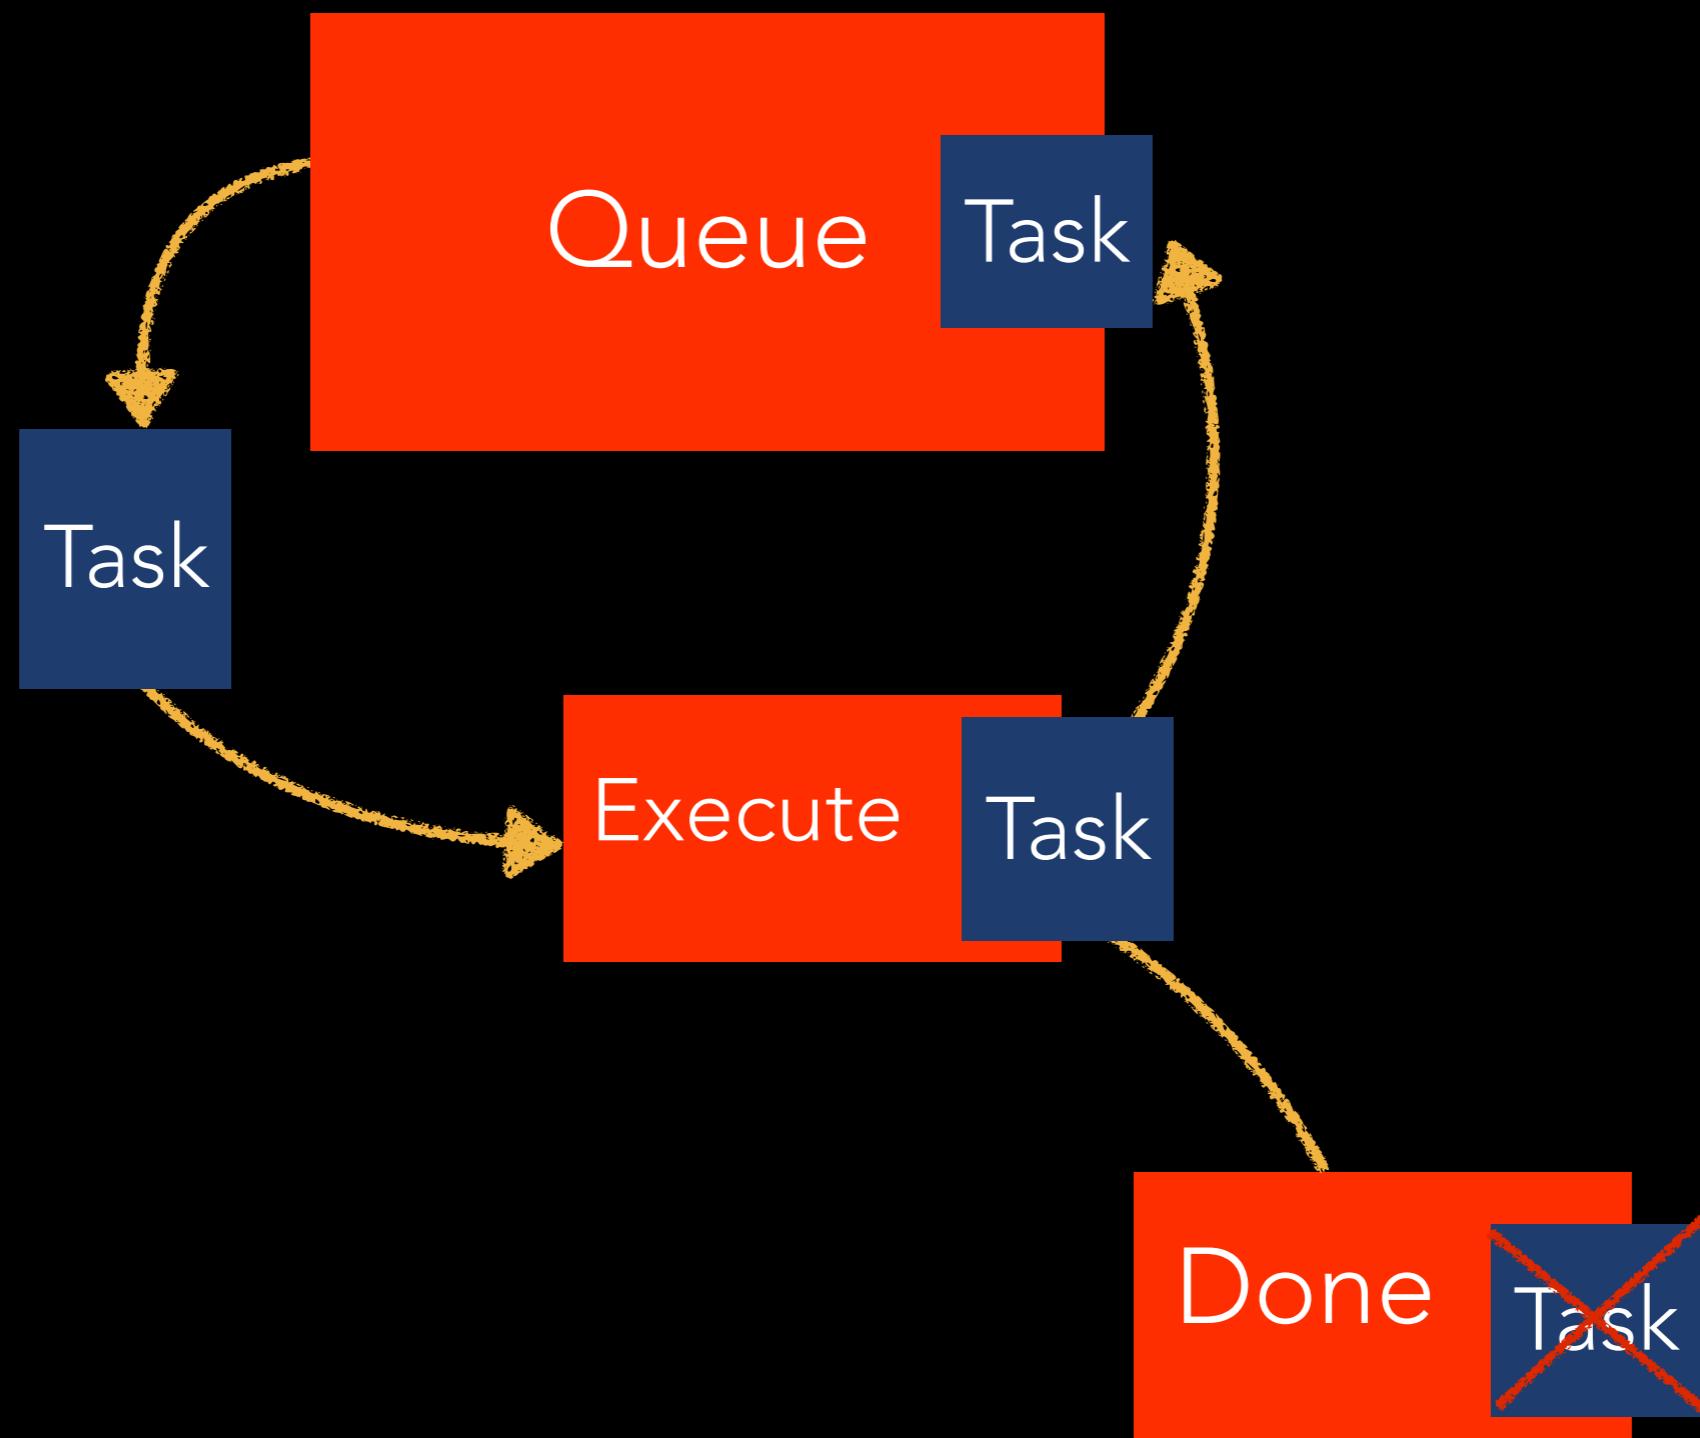
std::future<bool> TaskMaster::AddTrackedTask(std::unique_ptr<ITask> task, ITask* /*parent*/)
{
    assert(task);
    std::unique_ptr<TrackedTask> wrapper = std::make_unique<TrackedTask>(task);
    std::future<bool> status = wrapper->GetFuture();
    AddTask(std::move(wrapper));
    return status;
}
```

TASKSYSTEM (CONT) - PROCESSING A TASK

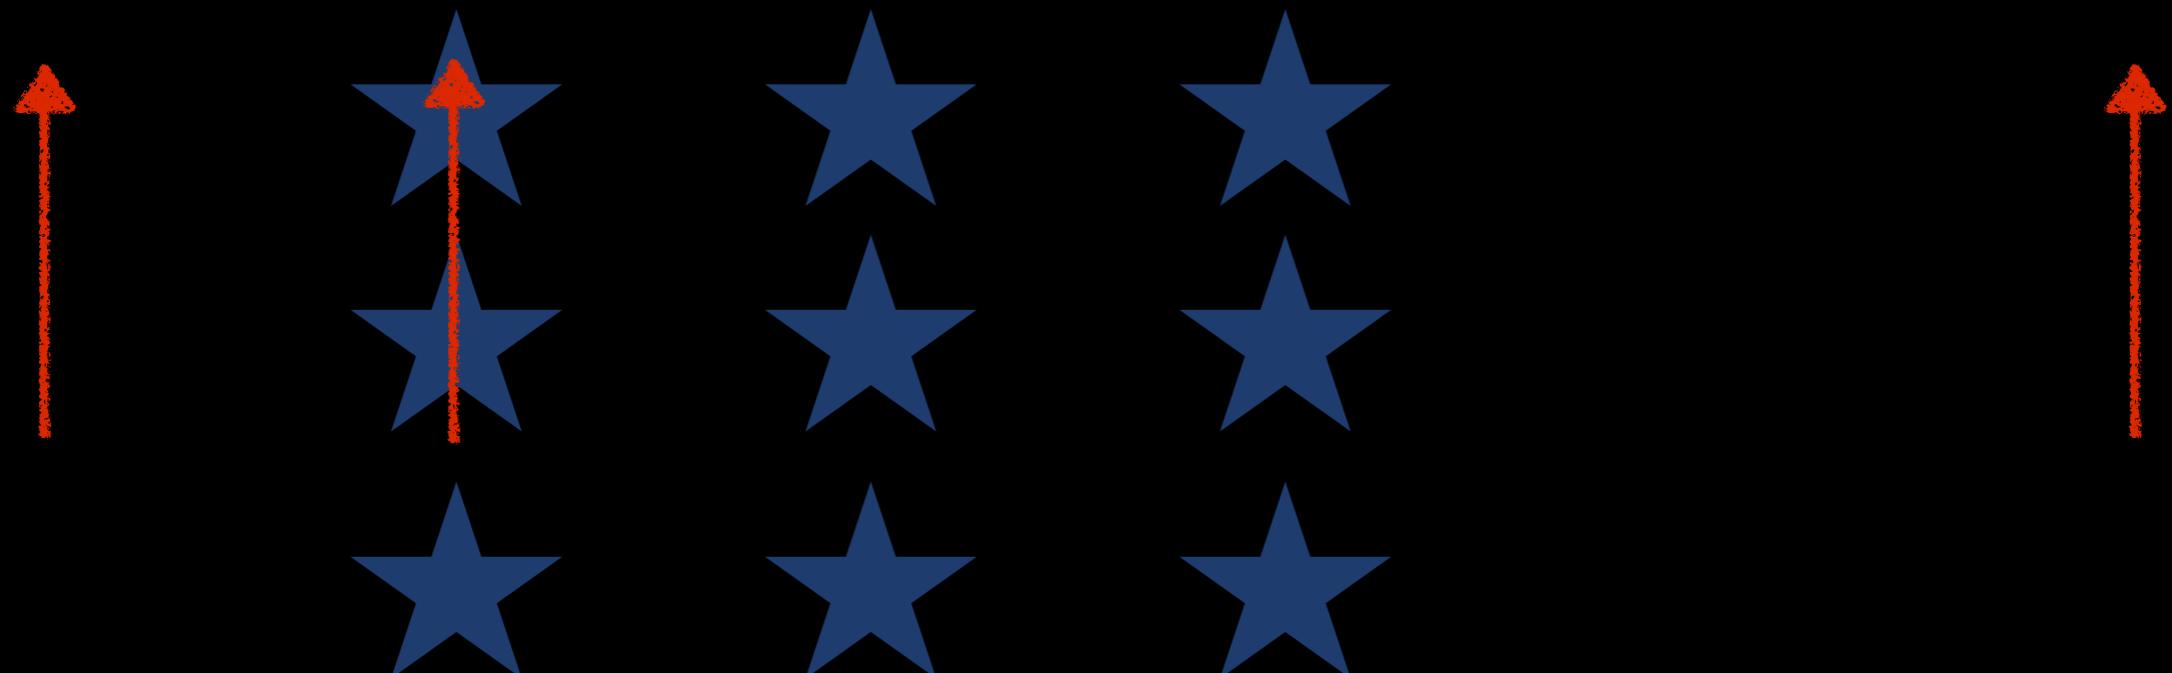
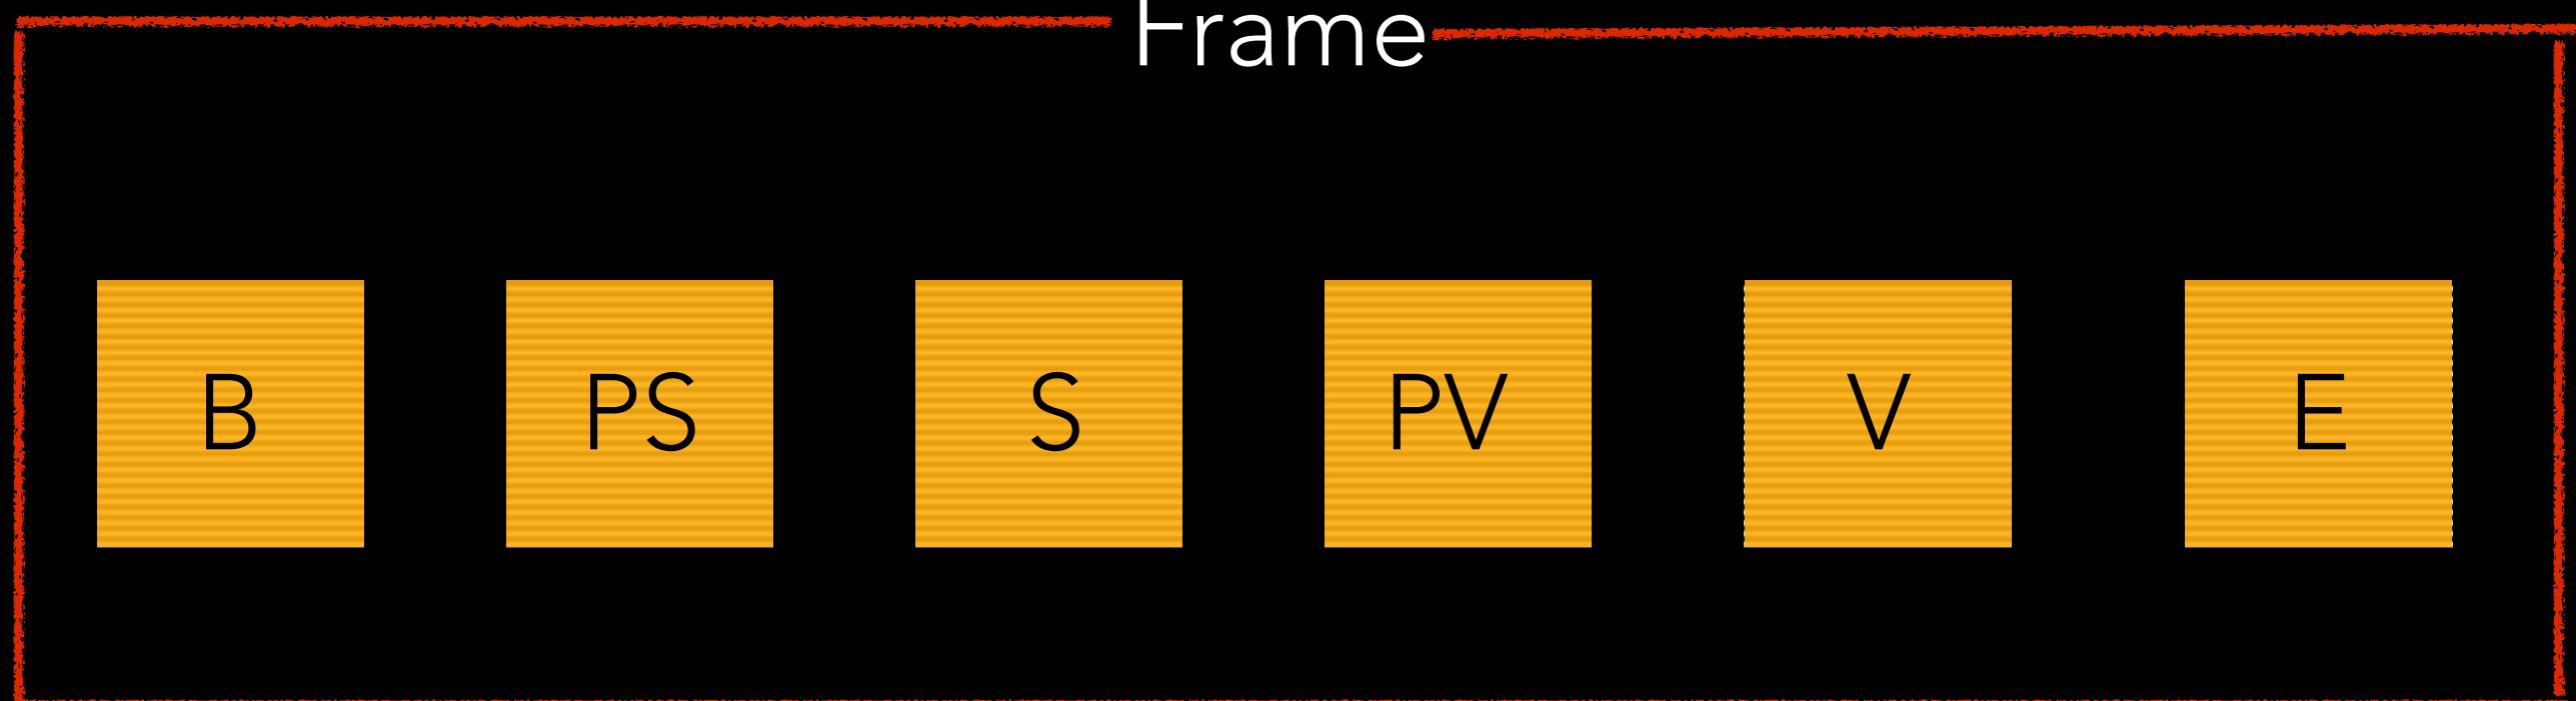
```
//Use std::thread::hardware_concurrency() to populate thread "pool" during initialize
std::atomic_bool s_RunThreads;
std::vector<std::unique_ptr<std::thread>> s_ThreadPool;

void runTask()
{
    std::thread::id id = std::this_thread::get_id();
    while (s_RunThreads)
    {
        std::unique_ptr<ITask> task = s_Tasks.pop(); ←
        if (task != nullptr)
        {
            if (!s_RunThreads)
            {
                //no longer need to keep the task around....we are shutting down
                break;
            }
            task->Run(); //Run it ←
            if (!task->IsComplete())
            {
                if (!s_RunThreads)
                {
                    //no longer need to keep the task around....we are shutting down
                    break;
                }
                s_Tasks.push(std::move(task)); ←
            }
        }
        else
        {
            std::this_thread::sleep_for(std::chrono::milliseconds(3));
        }
    }
}
```

TASKSYSTEM (CONT)-TASK FLOW



WHAT'S IN A FRAME?



WHAT'S IN A FRAME (CONT)

- Why is visualize not using parallelism?
 - prep for visualization does all prep-work of sorting and filtering visualize requests.
 - Visualize uses the gpu which typically is not multi threaded
 - SIMD can be used to help with processing data, but most prep-work should be done in prep stage.
 - Visual fancy features can be expensive so filtering and sorting should be done beforehand.

C++ STANDARD HIGHLIGHTS

- Standardization helps simplify code bases.
 - Move semantics speeding up data copying for all std data containers
 - Threading library
 - Algorithms Library
 - Algorithms work on raw arrays as well.
 - `std::distance` (getting index info from contiguous array members)

C++ STANDARD HIGHLIGHTS CONTINUED

- `Std::function`
- `static_assert`
- `Constexpr`
- `Std::tuple` (dynamic data grouping)
- Type traits and compile time checks
- `Std::atomics`
- `Std::array`, `std::unordered_map`, `std::unordered_set`

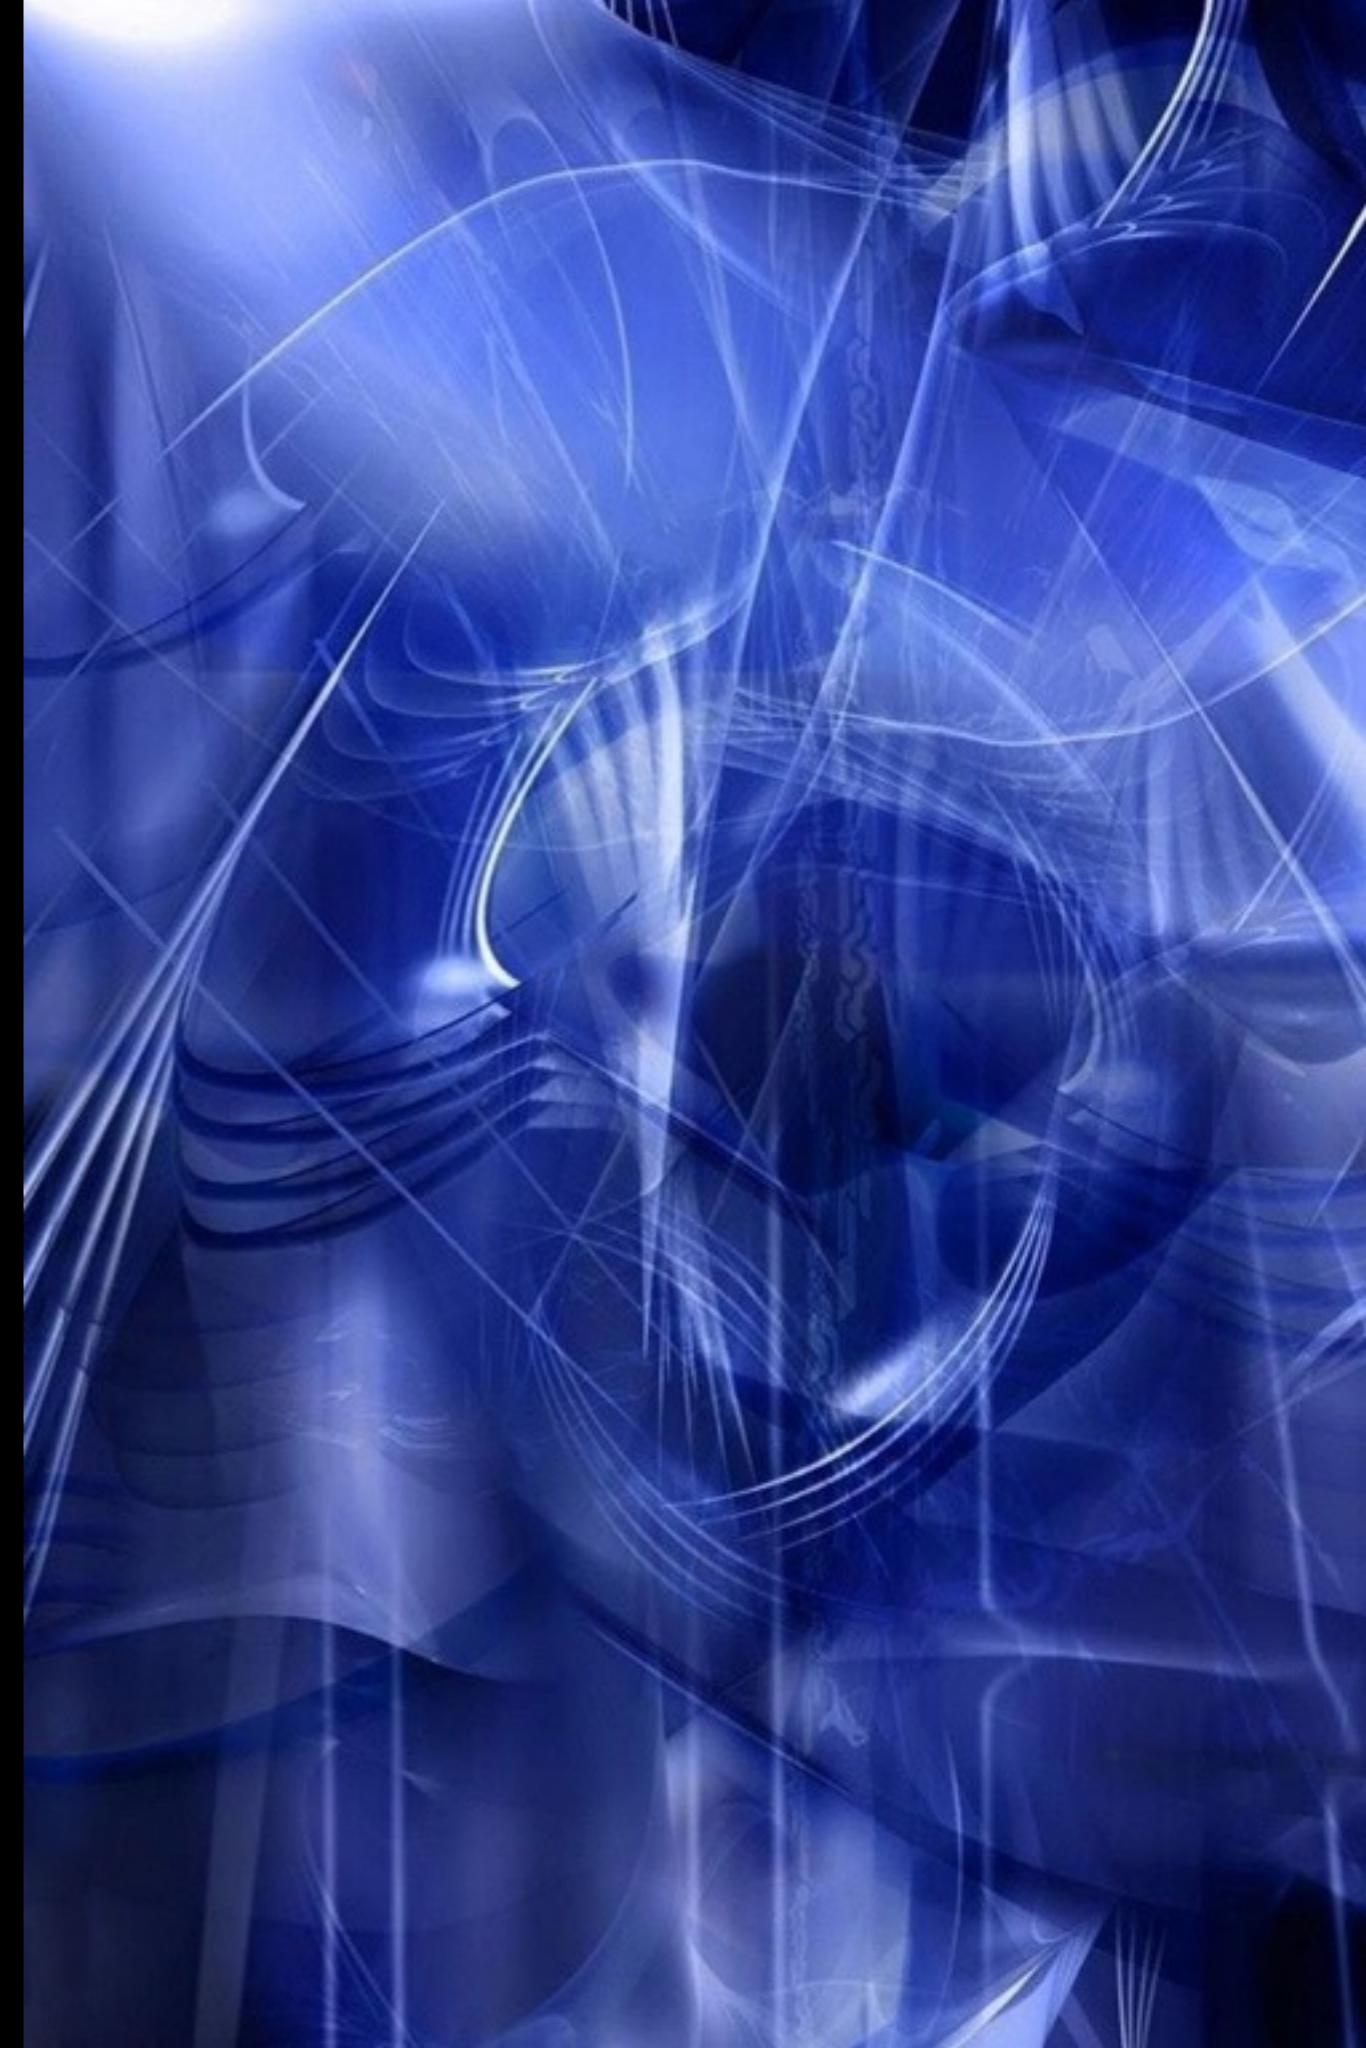


RIDDLE:

What can hold onto something bigger than itself and ensure proper clean-up?

C++ STANDARD HIGHLIGHTS SMART_PTRS

- Between class/struct data management
- clean up allocated memory (on proper shutdown and improper shutdown)
- Show how an object is supposed to be used
 - unique_ptr I am the owner (i.e. tasks, managers)
 - shared_ptr Ref counted item (i.e. asset resources)



RIDDLE

I am in front of you, you cannot see me, but you
may see me later?

C++ STANDARD HIGHLIGHTS

STD::FUTURE

- A tool for getting results from tasks
- Provide dependency link between tasks
- The size of a std::future is equal to size of 2 pointers

X86	X64
8	16

NOTE: `sizeof(std::future<bool>)` == `sizeof(std::future<int*>)`

C++ STANDARD HIGHLIGHTS

STD::FUTURE (CONT)

```
template<typename t>
bool futureReady(std::future<t>& toCheck)
{
    return toCheck.valid() && toCheck._Is_ready(); //NOTE: MS specific...
}

std::future<std::unique_ptr<SystemActionGather::Result>> m_SystemActionGatherResult;
std::future<bool> m_SystemActionApply;

void PrepForSimulate::Execute()
{
    //collate all of the previous frames actions.
    if (!m_triggeredGather)
    {
        std::unique_ptr<SystemActionGather> temp = std::make_unique<SystemActionGather>(m_FrameContext);
        m_SystemActionGatherResult = temp->GetResultFuture(); ←
        TASK_MASTER::AddTask(std::move(temp));
        m_triggeredGather = true;
    }
    else if (!m_triggeredApply && futureReady(m_SystemActionGatherResult)) ←
    {
        m_SystemActionApply = TASK_MASTER::AddTrackedTask(std::move(std::make_unique<SystemActionApply>(m_FrameContext, m_SystemActionGatherResult.get()))); ←
        m_triggeredApply = true;
    }
    else if (m_triggeredGather && m_triggeredApply && futureReady(m_SystemActionApply))
    {
        //FrameContext state data should now be const.... from here on out.
        TASK_MASTER::AddTask(std::move(std::make_unique<Simulate>(m_FrameContext)));
        return; // no longer loop
    }
    WantToExecuteAgain();
}
```

RIDDLE

What is infinite in size, but you never seem to have enough of?

C++ STANDARD HIGHLIGHTS

STD::CHRONO

- Very important for animations and logic timing
- Typically you have to be very careful of units in use
 - `unsigned int m_TimeStamp; // is this ms? s? ns?`
 - `unsigned int m_TimeSeconds; // better`
 - `m_TimeSeconds = 1000 //whoops did you mean 1 second or 1000ms or 1000s`
- Chrono's strongly typed classes are very helpful in knowing context of time
 - Chrono makes the conversions between units obvious and will fail to compile if conversion is invalid

C++ STANDARD HIGHLIGHTS

STD::CHRONO

```
std::chrono::high_resolution_clock::time_point now = std::chrono::high_resolution_clock::now();
std::chrono::milliseconds diff = std::chrono::duration_cast<std::chrono::milliseconds>(now - internal::s_LastFrame);
found->m_TimeStep = std::min(std::chrono::milliseconds(15), diff);
internal::s_LastFrame = now;
```

```
std::chrono::seconds test(1);
std::chrono::milliseconds testms(500);

std::chrono::milliseconds result = test + testms;
```

```
1s
1500ms
1500000000ns
```

```
std::cout << "..." << std::chrono::duration_cast<std::chrono::seconds>(result).count() << "s\n..."
→ << std::chrono::duration_cast<std::chrono::milliseconds>(result).count() << "ms\n..."
→ << std::chrono::duration_cast<std::chrono::nanoseconds>(result).count() << "ns\n";
```

C++ STANDARD HIGHLIGHTS

STD::CHRONO (CONT)

```
std::chrono::seconds test(1);
std::chrono::milliseconds testms(500);

std::chrono::seconds result = test + testms;
```

std::chrono::seconds test

no suitable user-defined conversion from "std::chrono::duration<long long, std::milli>" to "std::chrono::seconds" exists

```
std::chrono::seconds test(1);
```

1s
1000ms
1000000000ns

```
std::cout << "..." << test.count() << "s\n..."
```



```
→ << std::chrono::duration_cast<std::chrono::milliseconds>(test).count() << "ms\n..."
```

```
→ << std::chrono::duration_cast<std::chrono::nanoseconds>(test).count() << "ns\n";
```

WHAT ABOUT STD::ASYNC?

- Could spin up a new thread for each call.
- Could use same thread to execute.
- A bit more wild west than is desired.

LANGUAGE FEATURES ON THE WAY

- Filesystem Library (C++17) — Available as experimental now
- std::variant (C++17)
 - resource sterilization/de-serialization
 - As a statemachine



LANGUAGE FEATURES ON THE WAY (CONT)

- `Std::optional` (C++17)
 - Removed usage of sentinel objects
 - `bool m_triggeredSomething`
 - `Std::unique_ptr<Result> m_resultFromSomething`
 - Becomes
 - `std::optional< std::unique_ptr<Result> > m_resultFromSomething`
- Various Refinements

LANGUAGE NICE TO HAVE THINGS

- Pool alloc/dealloc of objects (lock free pool)
- “Pure” keyword to make pure virtual functions in class/struct
 - Virtual void Function () =0; ==> pure void Function (); or void Function () =pure;
 - P0078R0: <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2015/p0078r0.pdf>
- Standardized Network Lib
 - N4332: Networking Library Proposal (<https://isocpp.org/blog/2014/11/n4332>)

LANGUAGE NICE TO HAVE THINGS CONT

- Ability to set a name/affinity to a thread through std ... (if not using async/coroutines)
- std::atomic< std::unique_ptr<Object> >
 - N4058: Atomic Smart Pointers (<https://isocpp.org/blog/2014/06/n4058>)
- std::future<>.is_ready() ... (MS has one as part of their implementation _Is_ready)?
- std::future.valid() is not as clear as expected.
 - Valid does not mean connected to promise or otherwise initialized (So use of sentinel is required. Also after get ... state seems undefined)

LANGUAGE NICE TO HAVE THINGS CONT

- `constexpr` does not always mean compile time executed `constexpr` functions
 - `const int test = constexpr_hash("TEST"); // This is executed at runtime`
 - `constexpr const int test = constexpr_hash("TEST"); // This is executed at compile time`
- would like
 - `const int test = constexpr_hash("TEST"); // This is executed at compile time`
 - `const int test = 12345678; //after compile time executes`
 - `constexpr const int test = constexpr_hash("TEST"); // This is executed the same`

LANGUAGE NICE TO HAVE THINGS CONT

- smart_ptrs that can handle forward declarations with default destruction.
- class ClassName; (implementation in ClassName.h/.cpp) using default destruction
- std::unique_ptr<ClassName> m_Ptr; //fails to compile



REFERENCES

- www.cppreference.com
- www.isocpp.org
- GDC 2015: Christian Gyrling Parallelizing the Naughty Dog Engine Using Fibers
 - <http://www.gdcvault.com/play/1022186/Parallelizing-the-Naughty-Dog-Engine>
 - http://twvideo01.ubm-us.net/o1/vault/gdc2015/presentations/Gyrling_Christian_Parallelizing_The_Naughty.pdf
- C++ Seasoning Sean Parent <https://www.youtube.com/watch?v=qH6sSOr-yk8>
- CppCon 2015: Sean Parent "Better Code: Data Structures" <https://www.youtube.com/watch?v=sWgDk-o-6ZE>
- CppCon 2014: Mike Acton "Data-Oriented Design and C++" <https://www.youtube.com/watch?v=rX0ltVEVjHc>
- CppCon 2014: Nicolas Fleury "C++ in Huge AAA Games" <https://www.youtube.com/watch?v=qYN6eduU06s>
- Modern C++: What You Need to Know <https://www.youtube.com/watch?v=TJHgp1ugKGM>
- CppCon 2015: Nicolas Guillemot & Sean Middleditch "Birth of Study Group 14..." <https://www.youtube.com/watch?v=0tFxdreVx6s>
- CppCon 2015: Michael Wong "C++11/14/17 atomics and memory model..." <https://www.youtube.com/watch?v=DS2m7T6NKZQ>
- CppCon 2015: Fedor Pikus PART 2 "Live Lock-Free or Deadlock (Practical Lock-free Programming)" <https://www.youtube.com/watch?v=1obZeHnAwz4>
- CppCon 2014: Herb Sutter "Back to the Basics! Essentials of Modern C++ Style" <https://www.youtube.com/watch?v=xnqTKD8uD64>

Summary

- Standardized code is easier to maintain
- Standardized code is broadly tested
- C++11 offers many useful features for games
 - std::thread
 - smart pointers
 - std::chrono (timing lib)
 - std::future, std::atomic
 - constexpr

jasonjurecka@gmail.com / jjurecka@blizzard.com