

## Leak-Freedom in C++... *By Default*

Herb Sutter

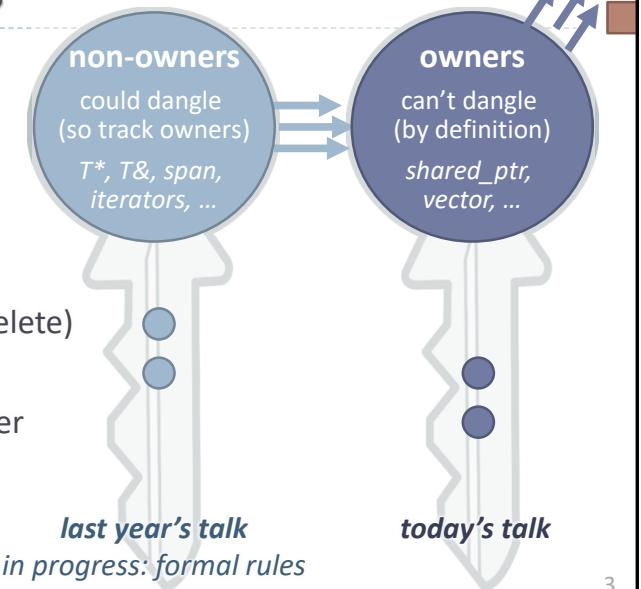
## Lifetime guarantees

Lifetime safety targets  
these classes of bugs:

**no dangling/invalid deref** (use after delete)

**no null deref**

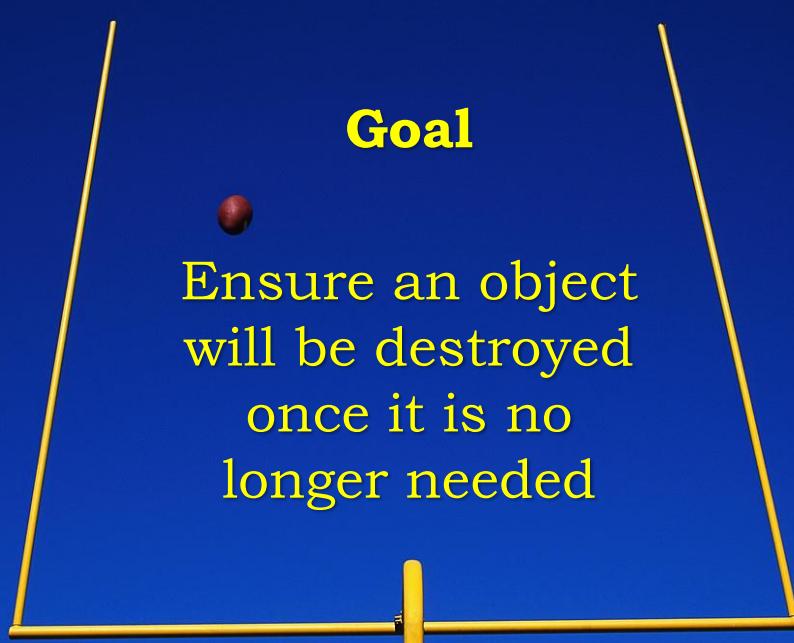
**no leaks** (delete objects when no longer  
used, and delete only once)



3

## Goal

Ensure an object  
will be destroyed  
once it is no  
longer needed



## “Leak freedom in C++” poster

Strategy	Natural examples	Cost	Rough frequency
<b>1. Prefer scoped lifetime by default</b> (locals, members)	Local and member objects – directly owned	Zero: Tied directly to another lifetime	O(80%) of objects
<b>2. Else prefer make_unique &amp; unique_ptr or a container</b> , if the object must have its own lifetime (i.e., heap) and ownership can be unique w/o owning cycles	Implementations of trees, lists	Same as new/delete & malloc/free <b>Automates</b> simple heap use in a library	O(20%) of objects
<b>3. Else prefer make_shared &amp; shared_ptr</b> , if the object must have its own lifetime (i.e., heap) and shared ownership w/o owning cycles	Node-based DAGs, incl. trees that share out references	Same as manual reference counting (RC) <b>Automates</b> shared object use in a library	

**Don't use owning raw \*'s == don't use explicit delete**

**Don't create ownership cycles** across modules by owning “upward” (violates layering)

Use `weak_ptr` to break cycles

5

## Example: HAS-A

- ▶ **Q:** What's the natural ownership abstraction for containing another object?
  - ▶ Part of my class's representation.
- ▶ **A: Nonstatic data member.**
  - ▶ Surprise! (not really)

```
class MyClass {
    Data data;
    /*...*/
};
```



6

## Example: HAS-A

- ▶ **Q:** What's the natural ownership abstraction?
- ▶ Part of my class's representation.
- ▶ **A: Nonstatic data member.**
- ▶ Surprise! (not really)

**Key:** Declares lifetime by construction.

**Correct:** Can see without looking at function bodies that there are no leaks.

**Efficient:** Zero extra allocation/tracking.

**Robust:** Strictly nested lifetime. Any bugs are sins of commission, not of omission.

```
class MyClass {
    Data data;
    /*...*/
};
```

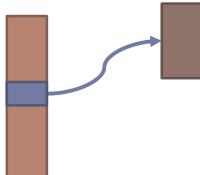


7

## Example: Decoupled HAS-A

- ▶ **Q:** What's the natural ownership abstraction for a decoupled member?
- ▶ Optional part (e.g., lazily initialized on demand).
- ▶ Changeable part (e.g., could change derived type).
- ▶ **A: unique\_ptr.**
- ▶ Note: If it shouldn't be null, remember to write your own move operations appropriately.

```
class MyClass {
    unique_ptr<Data> pdata;
    /*...*/
};
```



8

## Example: Decoupled

- ▶ **Q:** What's the natural ownership abstraction?
- ▶ Optional part (e.g., lazily initialized on demand)
- ▶ Changeable part (e.g., could change dynamically)
- ▶ **A: unique\_ptr.**
- ▶ Note: If it shouldn't be null, remember to write your own move operations appropriately.

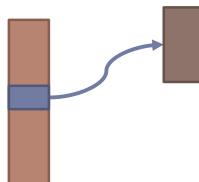
```
class MyClass {
    unique_ptr<Data> pdata;
    /*...*/
};
```

**Key:** Declares lifetime by construction.

**Correct:** Can see without looking at function bodies that there are no leaks.

**Efficient:** Equal space & time to correctly written manual new/delete by hand.

**Robust:** Strictly nested lifetime.  
(But remember to prevent null-after-move.)



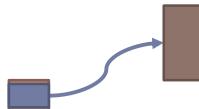
9

## Example: Pimpl Idiom (compilation firewall)

- ▶ **Q:** What's the natural ownership abstraction for a compilation firewall (Pimpl)?
- ▶ **A: const unique\_ptr.**
- ▶ Non-const can work too, but is more brittle because default move semantics are probably incorrect.
- ▶ “Another/detached part of ‘this’ object.”

```
template<class T>
using Pimpl = const unique_ptr<T>;
```

```
class MyClass {
    class Impl;      // defined in .cpp
    Pimpl<Impl> pimpl;
    /*... Note: declare destructor and write it elsewhere ...*/
};
```



10

## Example: Pimpl Idiom

- ▶ **Q:** What's the natural ownership abstraction?
- ▶ **A: const unique\_ptr.**
  - ▶ Non-const can work too, but is more brittle because default move semantics are probably incorrect.
  - ▶ “Another/detached part of ‘this’ object.”

**Key:** Declares lifetime by construction.

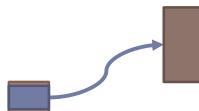
**Correct:** Can see without looking at function bodies that there are no leaks.

**Efficient:** Equal space & time to correctly written manual new/delete by hand.

**Robust:** Strictly nested lifetime. Any bugs are sins of commission, not of omission.

```
template<class T>
using Pimpl = const unique_ptr<T>;
```

```
class MyClass {
    class Impl;      // defined in .cpp
    Pimpl<Impl> pimpl;
    /*... Note: declare destructor and write it elsewhere ...*/
};
```

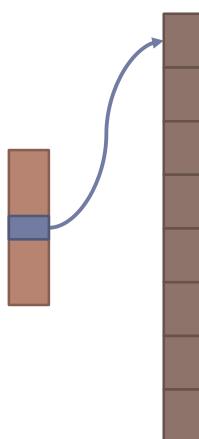


11

## Example: Dynamic array member

- ▶ **Q:** How can we express a fixed-but-dynamic-size member array?
- ▶ **A: const unique\_ptr<[]>.**
  - ▶ “Another/detached part of ‘this’ object.”
  - ▶ (Can remove *const* if you want whole-array move and handling for nullness.)

```
class MyClass {
    const unique_ptr<Data[]> array;
    int array_size; // or: Data* end;
    /*...*/
    MyClass(size_t num_data)
        : array(make_unique<Data[]>(num_data))
    { }
};
```

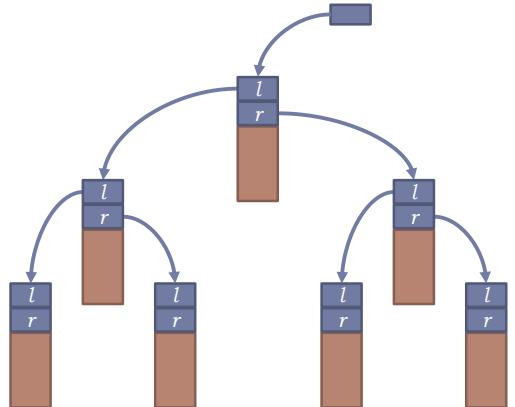


12

## Example: Tree (usually encapsulated)

- ▶ **Q:** What's the natural ownership abstraction for a tree?
- ▶ **A:** `unique_ptr`.

```
class Tree {
    struct Node {
        vector<unique_ptr<Node>> children;
        /* ... data ... */
    };
    unique_ptr<Node> root;
    /* ... */
};
```



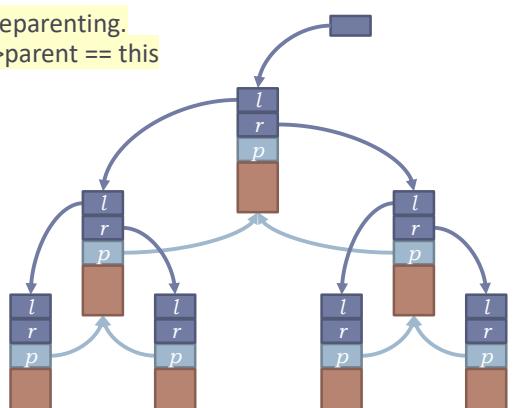
13

## Example: Tree (usually encapsulated)

- ▶ **Q:** What's the natural ownership abstraction for a tree?
- ▶ **A:** `unique_ptr`.

Only manual part: Update parent pointer when reparenting.  
Enforce invariant: `left->parent == this && right->parent == this`

```
class Tree {
    struct Node {
        vector<unique_ptr<Node>> children;
        Node* parent;
        /* ... data ... */
    };
    unique_ptr<Node> root;
    /* ... */
};
```



14

## Example: Tree (usual)

- ▶ **Q:** What's the natural ownership abstraction?
- ▶ **A: unique\_ptr.**
  - ▶ Only manual part: Update parent pointer when moving.
  - ▶ Enforce invariant: `left->parent == this && right->parent == this`

```
class Tree {
    struct Node {
        vector<unique_ptr<Node>> children;
        Node* parent;
        /* ... data ... */
    };
    unique_ptr<Node> root;
    /* ... */
};
```

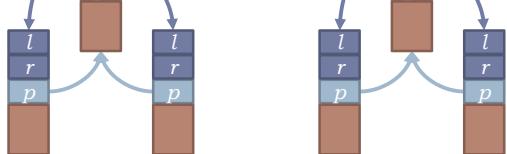
**Key:** Declares lifetime by construction.

**Correct:** Can see without looking at function bodies that there are no leaks.

**Efficient:** Equal space & time to correctly written manual new/delete by hand.

**Robust:** Any bugs are sins of commission, not omission. (Except for parent invariant.)

Exercise for the reader: Write the `reparent()` function that encapsulates updating the invariant. (Not hard, but more than two lines.)



15

Happiness is...

... a tree of *unique\_ptrs*



# Happiness is...



... not overflowing

## Releasing subtrees: Recursive vs. iterative

## Recursive (default)

- ▶ Destructors are recursive by default. Example:

## Releasing subtrees: Recursive vs. iterative

### Recursive (default)

- ▶ Destructors are recursive by default.  
Example:

```
void release_subtree(unique_ptr<Node> n)
{
} // destroy n here
// calls n->~Node()
// → n->children[0]->~Node();
// → n->children[0]->children[0]->~Node();
// → etc.
```

- ▶ Pro: Automatic and correct.
- ▶ Con: Unbounded stack depth.

### Iterative (manual options)

- ▶ Iterate by hand. Here's one example:

```
void release_subtree(unique_ptr<Node> n)
{
    while (n->children.size() > 0) {
        auto leaf = &n->children;
        while (leaf[0]->children.size() > 0)
            leaf = &leaf[0]->children;
        leaf.pop_front(); // found leaf, drop it
    }
} // destroy *n itself here
```

- ▶ Con: Manual optimization.
- ▶ Pro: Bounded stack depth.

▶ Note: This is  $O(N \log N)$ . There is an  $O(N)$  method but it doesn't fit on the slide.

19

## Releasing N-node subtree: Extra cost (on top of $O(N)$ dtors)

	+ Time	+ Stack space	+ Heap space
<b>Default:</b> prompt, recursive	—	$O(N)$	—
<b>Iterative, in-place:</b> prompt, navigate subtree and run dtors bottom-up (previous slide)	$+ O(N \log N)$	—	—
<b>Iterative, copy:</b> Move ptrs to the nodes to be pruned to a flattened local-scope heap list, then run dtors iteratively	$+ O(N)$	—	$O(N)$
<b>Iterative, deferred:</b> Move ptrs to the nodes to be pruned to a flattened side list, run destructors later iteratively	$+ O(N)$ <i>- dtors</i>	—	$O(N)$

Additional overhead, on top of running all the  $O(N)$  Node destructors themselves

“—” means minimal,  $O(K)$  with low  $K$

20

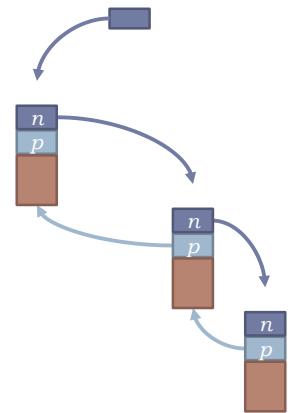
## Example: Doubly linked list

- ▶ **Q:** What's the natural ownership abstraction for a doubly linked list?

- ▶ **A: unique\_ptr.**

- ▶ Only manual part: Update previous pointer when relinking.  
Enforce invariant: `next->prev == this`

```
class LinkedList {
    struct Node {
        unique_ptr<Node> next;
        Node* prev;
        /* ... data ... */
    };
    unique_ptr<Node> root;
    /* ... */
};
```



21

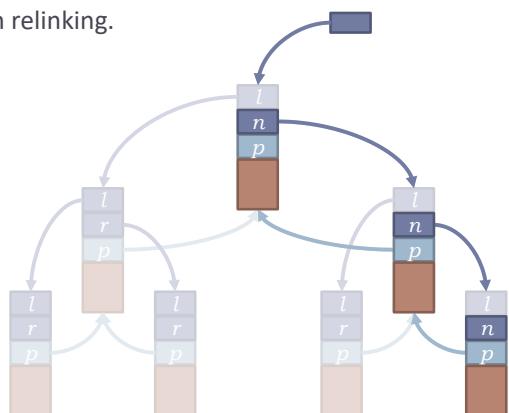
## Example: Doubly linked list

- ▶ **Q:** What's the natural ownership abstraction for a doubly linked list?

- ▶ **A: unique\_ptr.**

- ▶ Only manual part: Update previous pointer when relinking.  
Enforce invariant: `next->prev == this`

```
class LinkedList {
    struct Node {
        unique_ptr<Node> next;
        Node* prev;
        /* ... data ... */
    };
    unique_ptr<Node> root;
    /* ... */
};
```



22

## Example: Doubly linked list

- ▶ **Q:** What's the natural ownership abstraction for a doubly linked list?
- ▶ **A: unique\_ptr.**
  - ▶ Only manual part: Update previous pointer when inserting.
  - ▶ Enforce invariant: `next->prev == this`

```
class LinkedList {
    struct Node {
        unique_ptr<Node> next;
        Node* prev;
        /* ... data ... */
    };
    unique_ptr<Node> root;
    /* ... */
};
```

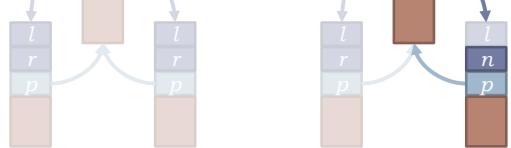
**Key:** Declares lifetime by construction.

**Correct:** Can see without looking at function bodies that there are no leaks.

**Efficient:** Equal space & time to correctly written manual new/delete by hand.

**Robust:** Any bugs are sins of commission, not omission. (Except for parent invariant.)

But, especially for the list case, be careful about destructor recursion. Typically, you want to prune iteratively.

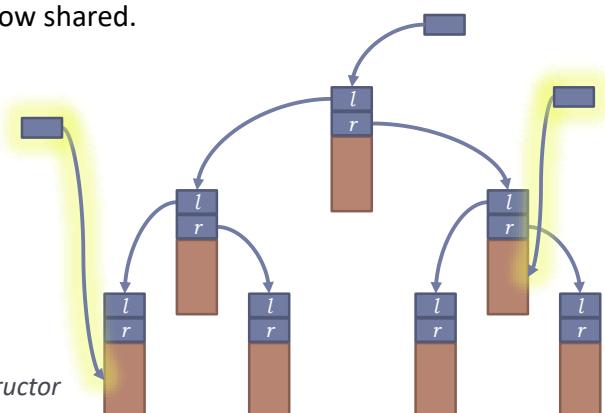


23

## Example: Tree that hands out strong refs

- ▶ **Q:** What's the natural ownership abstraction for a tree that hands out strong references to data?
- ▶ **A: shared\_ptr.** After all, ownership is now shared.

```
class Tree {
    struct Node {
        vector<shared_ptr<Node>> children;
        Data data;
    };
    shared_ptr<Node> root;
    shared_ptr<Data> find(/* ... */) {
        /* ... */ return {spn, &(spn->data)};
    } // note: uses shared_ptr aliasing constructor
};
```



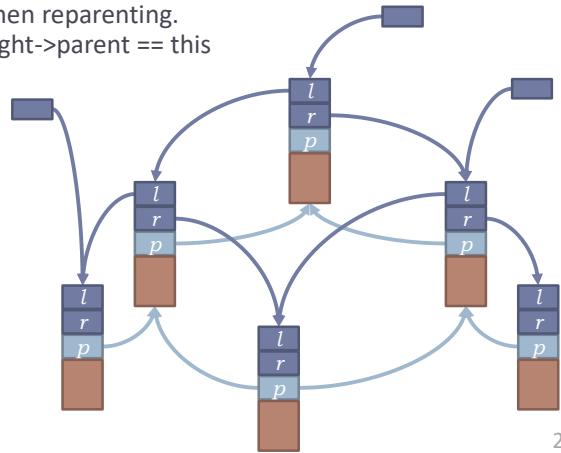
24

## Example: DAG of heap objects

- ▶ **Q:** What's the natural ownership for a node-based container that's a DAG?
- ▶ **A: shared\_ptr.**

- ▶ Only manual part: Update parent pointer when reparenting.  
Enforce invariant: left->parent == this && right->parent == this

```
class DAG {
    struct Node {
        vector<shared_ptr<Node>> children;
        vector<Node*> parents;
        /* ... data ... */
    };
    vector<shared_ptr<Node>> roots;
    /* ... */
};
```



25

## Example: DAG of heap objects

- ▶ **Q:** What's the natural ownership for a node-based container that's a DAG?
- ▶ **A: shared\_ptr.**

- ▶ Only manual part: Update parent pointer when reparenting.  
Enforce invariant: left->parent == this && right->parent == this

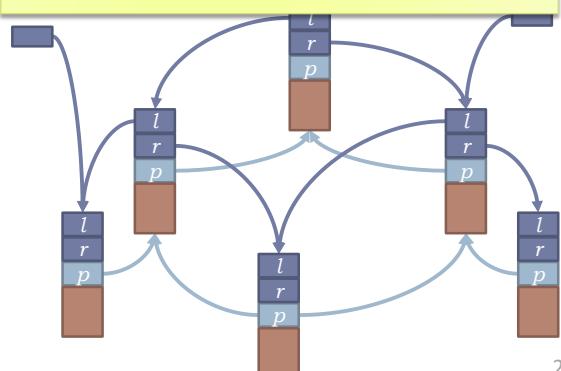
**Key:** Declares lifetime by construction.

**Correct:** Can see without looking at function bodies that there are no leaks.

**Efficient:** Equal space & time to correctly written manual new/delete + RC by hand.

**Robust:** Any bugs are sins of commission, not omission. (Except for parent invariant.)

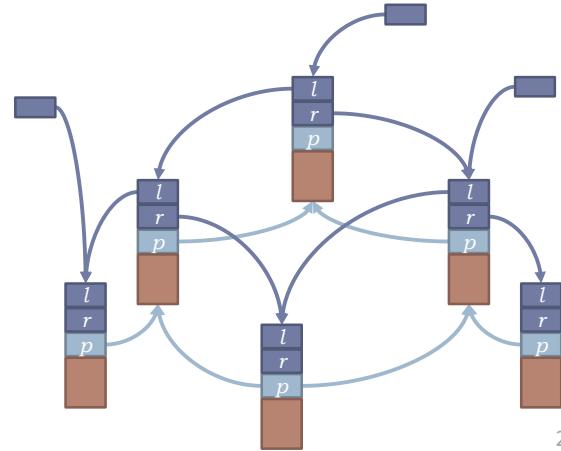
```
class DAG {
    struct Node {
        vector<shared_ptr<Node>> children;
        vector<Node*> parents;
        /* ... data ... */
    };
    vector<shared_ptr<Node>> roots;
    /* ... */
};
```



26

## Example: DAG of heap objs, unencapsulated

- ▶ **Q:** What's the natural ownership for objects that refer to each other w/o cycles?
- ▶ **A: shared\_ptr.**
  - Logically the same as the previous example.
  - Objects can be of different types.
  - Objects can come from different modules.
- ▶ Just don't violate layering...

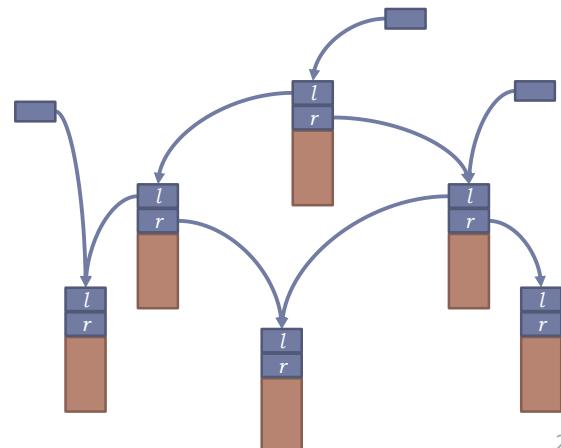


27

## Example: DAG of heap objs, unencapsulated

- ▶ **Q:** What's the natural ownership for objects that refer to each other w/o cycles?
- ▶ **A: shared\_ptr.**
  - Logically the same as the previous example.
  - Objects can be of different types.
  - Objects can come from different modules.

Might have parent pointers, or not.
- ▶ Just don't violate layering...

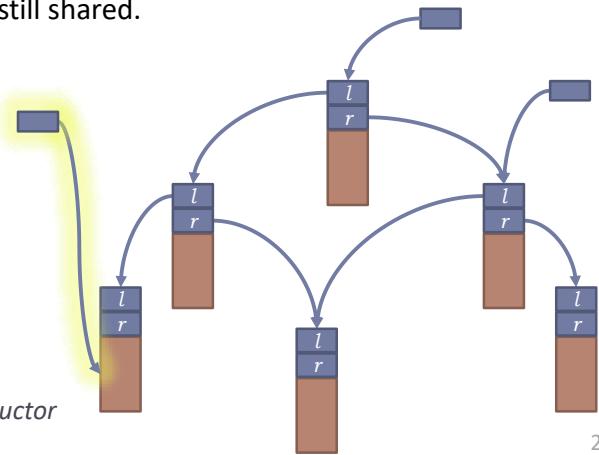


28

## Example: DAG that hands out strong refs

- ▶ **Q:** What's the natural ownership abstraction for a DAG that hands out strong references to data?
- ▶ **A: shared\_ptr.** As before, ownership is still shared.

```
class DAG {
    struct Node {
        vector<shared_ptr<Node>> children;
        vector<Node*> parents;
        /* ... data ... */
    };
    vector<shared_ptr<Node>> roots;
    shared_ptr<Data> find(/* ... */) {
        /* ... */ return {spn, &(spn->data)};
    } // note: uses shared_ptr aliasing constructor
};
```



29

## Example: Factory

- ▶ **Q:** What's the natural ownership when a factory returns a heap object?
- ▶ **A: unique\_ptr, or shared\_ptr + make\_shared.**
  - ▶ *unique\_ptr* if the object might not be shared by the code that then uses it.
  - ▶ *shared\_ptr + make\_shared* if the object will be shared by the code that then uses it.

`???_ptr<widget> make_widget( int id );`

```
unique_ptr<widget> make_widget( int id ) {
    return make_unique<widget>(id);
}
```

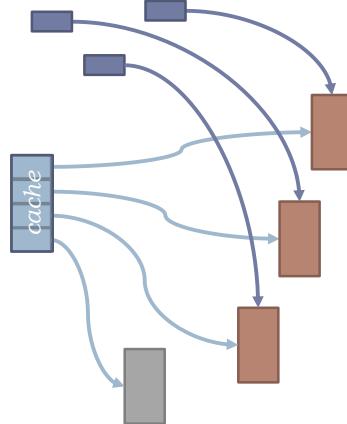
```
shared_ptr<widget> make_widget( int id ) {
    return make_shared<widget>(id);
}
```

30

## Example: Factory + cache

- ▶ **Q:** What's the natural ownership when a cache of heap objects doesn't keep objects alive (objects should live only as long as there are cache-external users)?
- ▶ **A: shared\_ptr + weak\_ptr.**

```
shared_ptr<widget> make_widget( int id ) {
    static map<int, weak_ptr<widget>> cache;
    static mutex mut_cache;
    lock_guard<mutex> hold( mut_cache );
    auto sp = cache[id].lock();
    if( !sp ) cache[id] = sp = load_widget( id );
    return sp;
}
```



31

## “Leak freedom in C++” poster

Strategy	Natural examples	Cost	Rough frequency
1. Prefer scoped lifetime by default (locals, members)	Local and member objects – directly owned	Zero: Tied directly to another lifetime	O(80%) of objects
2. Else prefer make_unique & unique_ptr or a container, if the object must have its own lifetime (i.e., heap) ownership can be unique w/o owning cycles	Implementations of trees, lists	Same as new/delete & malloc/free <b>Automates</b> simple heap use in a library	O(20%) of objects
3. Else prefer make_shared & shared_ptr, if the object must have its own lifetime (i.e., heap) and shared ownership w/o owning cycles	Node-based DAGs, incl. trees that share out references	Same as manual reference counting (RC) <b>Automates</b> shared object use in a library	

Don't use owning raw \*'s == don't use explicit delete

Don't create ownership cycles across modules by owning "upward" (violates layering)

Use weak\_ptr to break cycles

32

## “Leak freedom in C++” poster

Strategy	Natural examples	Cost	Rough frequency
1. Prefer scoped lifetime by default (locals, members)	Local and member objects – directly owned	Zero: Tied directly to another lifetime	O(80%) of objects
2. Else prefer <code>make_unique</code> & <code>unique_ptr</code> or a container, if the object must have its own lifetime (i.e., heap) ownership can be unique w/o owning cycles	Implementations of trees, lists	Same as new/delete & malloc/free <b>Automates</b> simple heap use in a library	O(20%) of objects
3. Else prefer <code>make_shared</code> & <code>shared_ptr</code> , if the object must have its own lifetime (i.e., heap) and shared ownership w/o owning cycles	Node-based DAGs, incl. trees that share out references	Same as manual reference counting (RC) <b>Automates</b> shared object use in a library	

Don't use owning raw \*'s == don't use explicit `delete` 



Don't create ownership cycles across modules by owning "upward" (violates layering)  
Use `weak_ptr` to break cycles

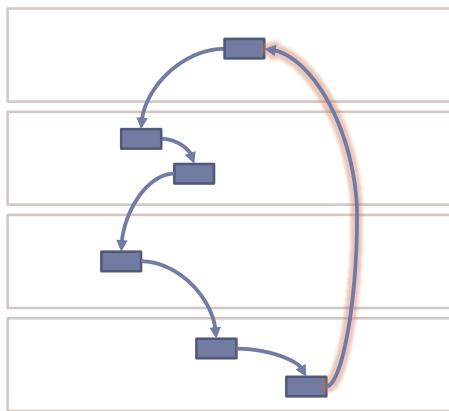


33

## Example: Cross-module cycle (prevent)

- ▶ **Q:** How can we avoid creating a cycle across modules?
- ▶ **A:** Don't own "upwards" (which violates layering).
  - ▶ How: Don't pass an owner down to "unknown code" that might store it.
  - ▶ Simple example: Storing a `shared_ptr`.

```
void bad(const shared_ptr<X>& x) {
    obj.register(x);
}
```

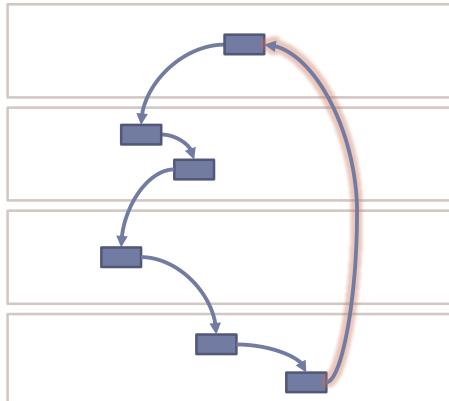


34

## Example: Cross-module cycle (prevent)

- ▶ **Q:** How can we avoid creating a cycle across modules?
- ▶ **A:** Don't own “upwards” (which violates layering).
  - ▶ How: Don't pass an owner down to “unknown code” that might store it.
  - ▶ Most commonly: Don't store an owner in a library **callback**.
    - ▶ NB: Direct parallel to “don't call unknown code while holding a lock” for concurrency.

```
void bad(const shared_ptr<X>& x) {
    obj.on_draw([=]{ x->extra_work(); });
}
```

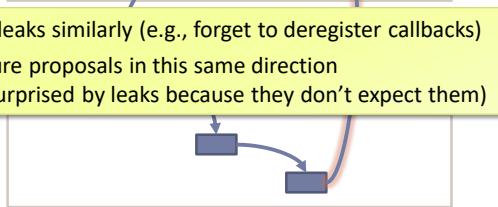
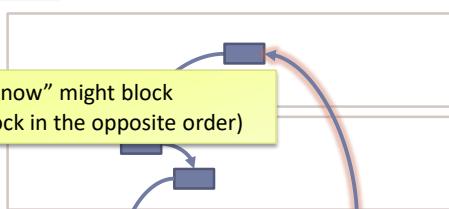


35

## Example: Cross-module cycle (prevent)

- ▶ **Q:** How can we avoid creating a cycle across modules?
- ▶ **A:** Don't own here, code that we “don't know” might store it...).
- ▶ How: Don't pass an owner down to “unknown code” that might store it.
- ▶ Most common ... and here, code that we “don't know” might block owner in a (in a way where another thread could block in the opposite order)
  - ▶ NB: Direct parallel to “don't call unknown code while holding a lock” for concurrency.

not “just a C++ problem” – Java/C# leaks similarly (e.g., forget to deregister callbacks)  
 ⇒ Java/C# have feature proposals in this same direction  
 void bad(const sha (their programmers are even more surprised by leaks because they don't expect them)  
 obj.on\_draw([=]{ x->extra\_work(); });
}

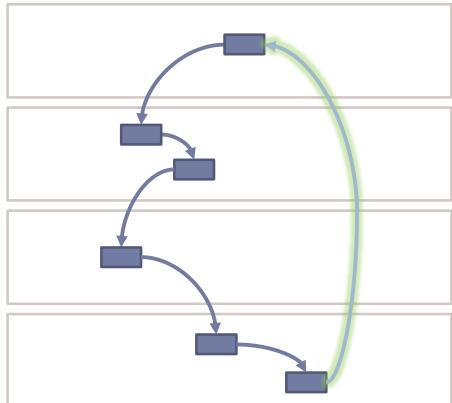


36

## Example: Cross-module cycle (prevent)

- ▶ **Q:** How can we avoid creating a cycle across modules?
- ▶ **A:** Don't own "upwards" (which violates layering).
  - ▶ How: Don't pass an owner down to "unknown code" that might store it.
  - ▶ Most commonly: Don't store an owner in a **callback**.
    - ▶ NB: Direct parallel to "don't call unknown code while holding a lock" for concurrency.

```
void good(const shared_ptr<X>& x) {
    obj.on_draw([w = weak_ptr<X>(x)]{
        if (auto x = w.lock()) x->extra_work(); });
}
```

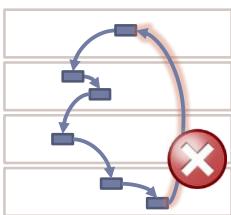


37

## Two categories of cycles

### Inter-module (compositional)

- ▶ Example: Composing two libs, each of which exposes access to heap objects.
  - ▶ Cycle is typically **heterogeneous**.
- ▶ **Answer: "Don't do that"** (just covered).
  - ▶ Don't violate layering by owning upward.
  - ▶ Don't store an owner in a callback.

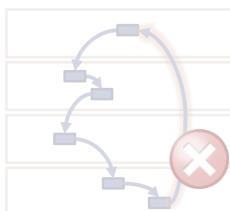


38

## Two categories of cycles

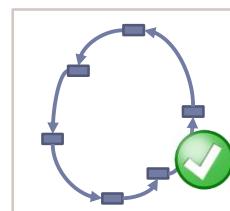
### Inter-module (compositional)

- Example: Composing two libs, each of which exposes access to heap objects.
  - Cycle is typically **heterogeneous**.
- Answer: “Don’t do that”** (just covered).
  - Don’t violate layering by owning upward.
  - Don’t store an owner in a callback.



### Intra-module (local)

- Example: Graph.
  - Cycle can be made **homogeneous**.
- Legitimate and reasonable.**
  - Not fully automated by `std::` facilities.
  - What should we do *today*? Could we maybe do better *tomorrow*?

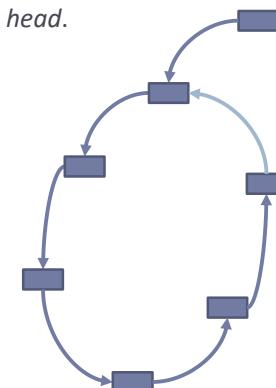


39

## Example: Circular list

- Q:** What’s the natural ownership abstraction for a circular list?
- A: Mostly `unique_ptr`** (in today’s portable C++).
  - Only manual part: `next()` to connect last/sentinel node to `head`.
  - Note: Easier because this is a **static cycle**.

```
class CircularList {
    class Node {
        unique_ptr<Node> next; // private
        unique_ptr<Node>& head;
    public:
        auto get_next() { return next ? next.get() : head.get(); }
        /*... ctors and data ...*/
    };
    unique_ptr<Node> head;
    /*...*/
};
```



40

## Example: Circular list

- ▶ **Q:** What's the natural ownership abstraction?
- ▶ **A: Mostly unique\_ptr** (in today's portable)
  - ▶ Only manual part: `next()` to connect last/second-to-last nodes.
  - ▶ Note: Easier because this is a **static cycle**.

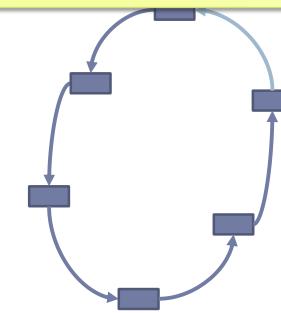
```
class CircularList {
    class Node {
        unique_ptr<Node> next; // private
        unique_ptr<Node>& head;
    public:
        auto get_next() { return next ? next.get() : head.get(); }
        /*... ctors and data ...*/
    };
    unique_ptr<Node> head;
    /*...*/
};
```

**Key:** Declares lifetime by construction.

**Correct:** Can see without looking at function bodies (except `next()`) that there are no leaks.

**Efficient:** Equal space & time to correctly written manual new/delete by hand.

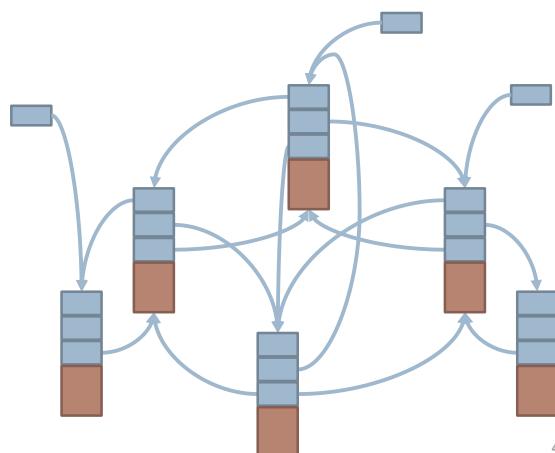
**Robust:** Any bugs are sins of commission, not omission. (Except for `next()` invariant.)



41

## Example: Graph of heap objs, unencapsulated

- ▶ **Q:** What's the natural ownership abstraction for a **possibly-cyclic** graph of objects?

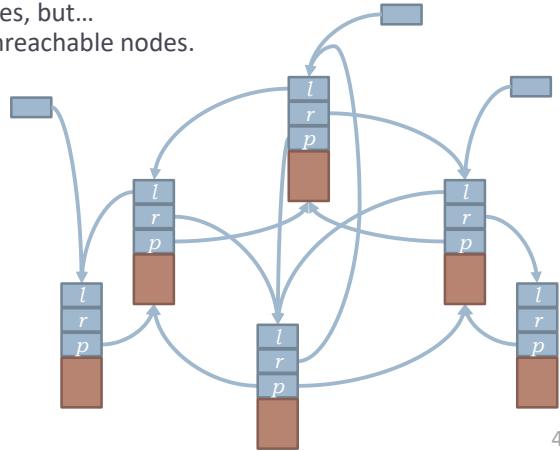


42

## Example: Graph of nodes, encapsulated

- ▶ **Q:** What's the natural ownership abstraction for a **possibly-cyclic** graph of nodes?
- ▶ **A: Partial** (in today's portable C++).
  - ▶ Cleanest solution is `vector<unique_ptr>` nodes, but...
    - ... **manual detection** is needed to identify unreachable nodes.

```
class Graph {
    struct Node {
        vector<Node*> children;
        vector<Node*> parents;
        /*... data ...*/
    };
    vector<Node*> roots;
    vector<unique_ptr<Node>> nodes; // ???
    /*...*/
};
```



43

## Example: Graph

- ▶ **Q:** What's the natural ownership abstraction for a possibly-cyclic graph of nodes?
- ▶ **A: Partial** (in today's portable C++).
  - ▶ Cleanest solution is `vector<unique_ptr>` nodes, but...
    - ... **manual detection** is needed to identify unreachable nodes.

```
class Graph {
    struct Node {
        vector<Node*> children;
        vector<Node*> parents;
        /*... data ...*/
    };
    vector<Node*> roots;
    vector<unique_ptr<Node>> nodes; // ???
    /*...*/
};
```

### Partly declares ownership by construction.

I know the nodes won't leak when the *Graph* is destroyed.

But I don't know without looking at function bodies that there are no leaks of *Nodes* during the lifetime of the *Graph*.

### Manually traversing the nodes to discover the unused ones is manual memory management.

Manually calling `nodes.erase(unreachable_node);` is functionally the same as manually calling `delete unreachable_node;`.

Our goal is **automatic** memory management.

*earlier this week  
at CppCon...*

*“How many of you know how to  
find a cycle in a directed graph?”*

— Chandler Carruth, Wednesday

in “Garbage In, Garbage Out: Arguing about  
Undefined Behavior with Nasal Demons”

Paging **Tobias Langner...**

```
set<Node*> GetConnectedNodes() {  
    set<Node*> result;  
    vector<Node*> stack; //manual stack to allow deeper nesting  
    stack.push_back(root); //start from root  
    while (!stack.empty()) { //continue until no work to do
```

Paging **Andre Kostur...**

```
set_difference(begin(allnodes), end(allnodes),
               begin(visited), end(visited),
               inserter(unreachable, begin(unreachable)));
```

Paging **Federico Lebron...**

```
void ShrinkToFit() {
    auto wahlberg = [this] { // get it? because... Mark.
```

## Example: Graph

**Challenge: Implement  
*Graph::remove\_unused\_nodes()***

- ▶ **Q:** What's the natural ownership abstraction for a possibly-cyclic graph of nodes?
- ▶ **A: Partial** (in today's portable C++).
  - ▶ Cleanest solution is `vector<unique_ptr>` nodes, but...
    - ... **manual detection** is needed to identify unreachable nodes.

```
class Graph {
    struct Node {
        vector<Node*> children;
        vector<Node*> parents;
        /*... data ...*/
    };
    vector<Node*> roots;
    vector<unique_ptr<Node>> nodes; // ???
    /*...*/
};
```

49

## Example: Graph

**Challenge: Implement  
*Graph::remove\_unused\_nodes()***

- ▶ **Q:** What's the natural ownership abstraction for a possibly-cyclic graph of nodes?
- ▶ **A: Partial** (in today's portable C++).
  - ▶ Cleanest solution is `vector<unique_ptr>` nodes, but...
    - ... **manual detection** is needed to identify unreachable nodes.

```
class Graph {
    struct Node {
        vector<Node*> children;
        vector<Node*> parents;
        /*... data ...*/
    };
    vector<Node*> roots;
    vector<unique_ptr<Node>> nodes; // ???
    /*...*/
};
```

```
void Graph::remove_unused_nodes() {
    vector<const Node*> orphans(nodes.size());
    transform(nodes.begin(), nodes.end(), orphans.begin(),
              [] (const auto& x){ return x.get(); });
    sort(orphans.begin(), orphans.end());

    for (const auto& node : all_nodes()) // or longer traversal code
        orphans.erase(lower_bound(
            orphans.begin(), orphans.end(), node.get()));

    for (const auto o : orphans)
        nodes.erase(find_if(nodes.begin(), nodes.end(),
                            [o](const auto& x){ return x.get() == o; }));
}
```

Pro: Probably have a traversal function anyway.  
 Con: "`nodes.erase`" is just as manual as "`delete`".  
 Con: Trimming requires full traversal + overhead.  
 Con: Pattern, needs to be handwritten each time.



Strategy	Natural examples	Cost	Rough frequency
1. Prefer scoped lifetime by default (locals, members)	Local and member objects – directly owned	Zero: Tied directly to another lifetime	O(80%) of objects
2. Else prefer <code>make_unique</code> & <code>unique_ptr</code> or a container, if the object must have its own lifetime (i.e., heap) and ownership can be unique w/o owning cycles	Implementations of trees, lists	Same as new/delete & malloc/free Automates simple heap use in a library	O(20%) of objects
3. Else prefer <code>make_shared</code> & <code>shared_ptr</code> , if the object must have its own lifetime (i.e., heap) and shared ownership w/o owning cycles	Node-based DAGs, incl. trees that share out references	Same as manual reference counting (RC) Automates RC heap use in a library	

Don't use owning raw \*'s == don't use explicit `delete`  
Don't create ownership cycles across modules by owning "upward" (violates layering)  
Use `weak_ptr` to break cycles

51

## Happiness is...

... scopes + *unique\_ptrs* + *shared\_ptrs*  
(and not too many cycles)

Questions?

## *std::\*ptr limitations*

- ▶ Ownership cycles  $\Rightarrow$  **leak** objects and memory.
  - ▶ *shared\_ptr* uses reference counting.
  - ▶ Plain reference counting leaks cycles.
  - ▶ Can't always be naturally broken using a *weak\_ptr*.
- ▶ Unbounded pointer copy cost  $\Rightarrow$  exceed **real-time** execution deadlines.
  - ▶ *shared\_ptr* calls destructors promptly. (Usually a great thing!)
  - ▶ But objects can own other objects transitively.
  - ▶ Plain assignment to a *shared\_ptr* can have arbitrarily high cost. (Similar issue with EH.)
- ▶ Unbounded destructor depth  $\Rightarrow$  exceed **stack limits** in constrained environments.
  - ▶ Scoped/*unique\_ptr*/*shared\_ptr* lifetime calls destructors recursively (nested).
  - ▶ As above, ownership trees can be deep.
  - ▶ Releasing a tree of objects requires having sufficient stack space.

Let's start here.



53

## Example: Graph

- ▶ **Q:** What's the natural ownership abstraction?
- ▶ **A: Partial** (in today's portable C++).
  - ▶ Cleanest solution is *vector<unique\_ptr>* nodes ... **manual detection** is needed to identify unused nodes.

```
class Graph {
    struct Node {
        vector<Node*> children;
        vector<Node*> parents;
        /*... data ...*/
    };
    vector<Node*> roots;
    vector<unique_ptr<Node>> nodes; // ???
    /*...*/
};
```

Challenge: Implement  
*Graph::remove\_unused\_nodes()*

Alternatives move space/time overhead around.  
Example: Add an age count to each node, ++ on each full traversal, then *remove\_unused\_nodes()* is a 1-pass algorithm to trim nodes for which  $age < curr\_age$ .

But it's **manual** memory management.  
We'd like not to write a custom sweep by hand in every situation where a cycle could occur.

**Q:** Can we automate any part of this lifetime as a reusable library, like we automated new/delete with *unique\_ptr* and RC with *shared\_ptr*?

## Releasing N-nodes

**Recall from earlier:**  
 Remember recursive (default, but can consume unbounded stack)  
 vs.  
 iterative (manual, but stack-friendly)?

**Iterative, deferred:** Move ptrs to the nodes to be pruned to a flattened side list, run destructors later iteratively

Extra cost (on top of  $O(N)$  dtors)

+ Time	+ Stack space	+ Heap space
—	$O(N)$	—
$+ O(N \log N)$	—	—
$+ O(N)$	—	$O(N)$
<b>+ <math>O(N)</math> - dtors</b>	—	$O(N)$

Additional overhead, on top of running all the  $O(N)$  Node destructors themselves

“—” means minimal,  $O(K)$  with low  $K$

55

# W A R N I N G

## EXPERIMENT IN PROGRESS



[github.com/hsutter/gcpp](https://github.com/hsutter/gcpp)

56

1 object  
1 owner  
*unique\_ptr*

1 object  
 $N$  owners  
*shared\_ptr*



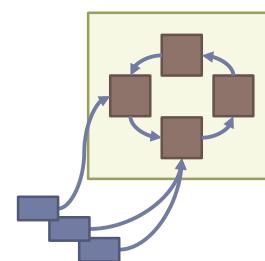
these work great when a single object's lifetime is manageable in isolation  
(e.g., just looking at immediate neighbors)

57

1 object  
1 owner  
*unique\_ptr*

1 object  
 $N$  owners  
*shared\_ptr*

**$N$  objects**  
1 or  $N$  owners  
*deferred\_ptr*  
(experimental)



these work great when a single object's lifetime is manageable in isolation  
(e.g., just looking at immediate neighbors)

reachability is a property of the whole group  
⇒ not detectable from 1 object or subgroup

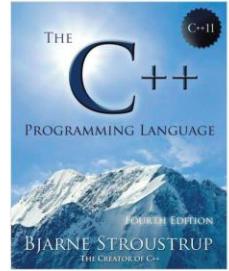
58

*earlier this week  
at CppCon...*

— Bjarne Stroustrup,  
Monday  
in “The Evolution of  
C++: Past, Present, and  
Future”

Morgan Stanley

## C++ in two lines



- Direct map to hardware
  - of instructions and fundamental data types
  - Initially from C
  - Future: use novel hardware better (caches, multicores, GPUs, FPGAs, SIMD, ...)
- Zero-overhead abstraction
  - Classes, inheritance, generic programming, ...
  - Initially from Simula (where it wasn't zero-overhead)
  - Future: Type- and resource-safety, concepts, modules, concurrency, ...



Stroustrup - CppCon'16

12

*earlier this week  
at CppCon...*

— Jason Turner,  
Wednesday  
in “Rich Code for Tiny  
Machines”

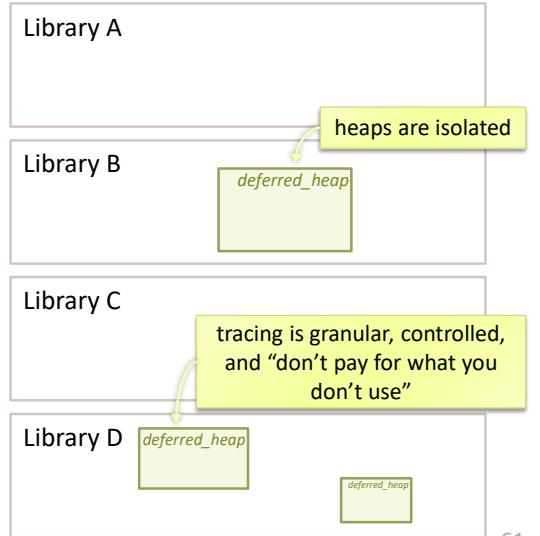
## ZERO OVERHEAD ABSTRACTIONS

- Objects
- Methods
- (Delegating) Constructors / Destructors
- Lambdas with/without captures
- Structured bindings (aka destructuring)
- Function calls that have been inlined
- If-init
- Templates (variadic, recursive)
- Standard algorithms

Jason Turner CppCon 2016 @lefticus  
<http://tinyurl.com/RichCodeCppCon2016>

## Overview: *deferred\_heap* + *deferred\_ptr*

- ▶ A *deferred\_heap* owns a “bubble” of objects that can point to each other arbitrarily.
  - ▶ Local heaps ⇒ isolation, composability.
  - ▶ Cycles must stay within one heap.
- ▶ *.make<T>()* creates an object, stores its destructor, and returns *deferred\_ptr<T>*.
- ▶ *.collect()* traces *this* heap, and destroys objects no longer reachable from roots.
  - ▶ Fully “don’t pay for what you don’t use.”
- ▶ *~deferred\_heap()* runs any pending destructors, nulls any *deferred\_ptr*s that outlive the heap, and releases its memory all at once like a region.



61

A working demo & source of ideas

Not production-quality (but feedback welcome)

Tries to automate things that require custom code today

⇒ another fallback “when needed,” not a primary option

## an experiment with **deferred** and **unordered** destruction

Goal is to automate tracing  
⇒ destructors run **later**, with  
**deterministic scope and timing**

Emphasis is on **destructors**  
⇒ **object lifetime**,  
not just memory lifetime  
⇒ no new “finalizer” concept

In destructors, pointers to other deferred objects are **null**  
⇒ safely unordered, enforces “no accidental use of a destroyed  
object” and “no resurrection”  
⇒ destructors of an object tree can be run **iteratively**, not nested

[github.com/hsutter/gcpp](https://github.com/hsutter/gcpp)

62

## Example: Graph

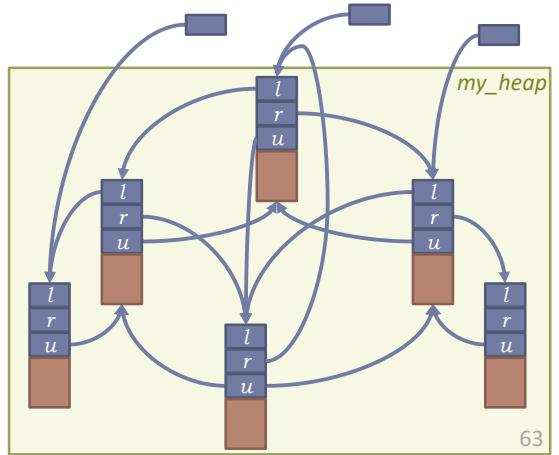
- ▶ **Q:** What's the natural ownership abstraction for a possibly-cyclic graph?

- ▶ **(experimental) A: deferred\_ptr.**

- ▶ No manual part needed:  
Can use `deferred_ptr` to the parent too.

```
class Graph {
    deferred_heap my_heap;
    struct Node {
        deferred_ptr<Node> left, right, up;
        /*... data ...*/
    };
    vector<deferred_ptr<Node>> roots;
    /*...*/
};
```

Private heap just for this object,  
modular & composable



63

## Example: Graph

- ▶ **Q:** What's the natural ownership abstraction

- ▶ **(experimental) A: deferred\_ptr.**

- ▶ No manual part needed:  
Can use `deferred_ptr` to the parent too.

```
class Graph {
    deferred_heap my_heap;
    struct Node {
        deferred_ptr<Node> left, right, up;
        /*... data ...*/
    };
    vector<deferred_ptr<Node>> roots;
    /*...*/
};
```

Private heap just for this object,  
modular & composable

Key: Declares lifetime by construction.

I know without looking at function bodies that there are no leaks. (I don't have to look at the body of `.collect()`, just call it when desired.)

```
void Graph::remove_unused_nodes() {
    heap.collect();
}
```

Pro: Automates tracing, doesn't need to be handwritten with custom code each time.

Pro: No manual "delete" masquerading as "`nodes.erase`".

Con: Trimming still requires full traversal + overhead (but limited to this heap ⇒ granular control).

54

54

## Example: DAG that hands out strong refs

- Q: What's the natural ownership abstraction for a DAG that hands out strong refs?

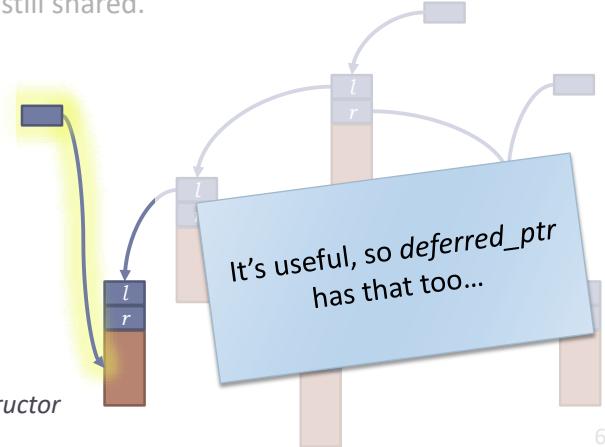
Recall from earlier:

Remember this example of the tree that handed out strong references, and the `shared_ptr` aliasing constructor?

```
class Node {
    vector<shared_ptr<Node>> children;
    /* ... data ... */
};

vector<shared_ptr<Node>> roots;
shared_ptr<Data> find(/*...*/) {
    /*...*/ return {spn, &(spn->data)};
} // note: uses shared_ptr aliasing constructor
};
```

The top is still shared.

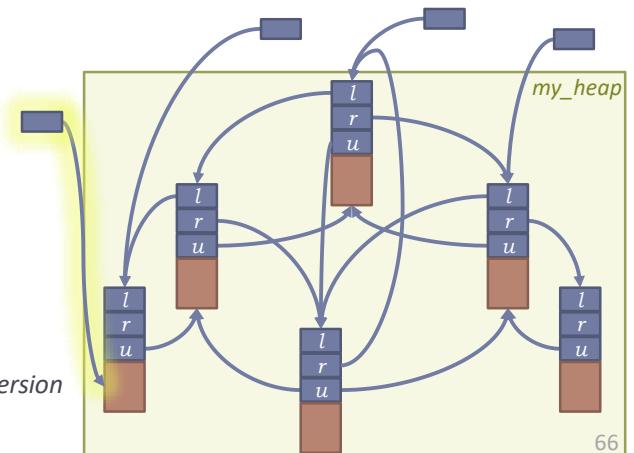


65

## Example: Graph that hands out strong refs

- Q: What's the natural ownership abstraction for a possibly-cyclic graph?
- (experimental) A: `deferred_ptr`.

```
class Graph {
    deferred_heap my_heap;
    struct Node {
        deferred_ptr<Node> left, right, up;
        /* ... data ... */
    };
    vector<deferred_ptr<Node>> roots;
    /* ... */
    deferred_ptr<Data> find(/*...*/) {
        /*...*/ return dpn.ptr_to(&Node::data);
    } // note: uses deferred_ptr aliasing conversion
};
```



66

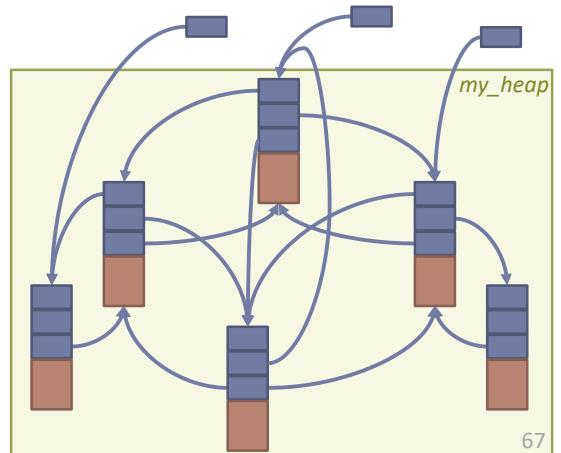
## Example: N-ary Graph

- ▶ **Q:** What if you want an N-ary Graph (not just left and right children)?

- ▶ **(experimental) A:** `vector<T, deferred_allocator<T>>`

- ▶ Convenience alias: `deferred_vector<T>`

```
class Graph {
    static deferred_heap my_heap;
    struct Node {
        deferred_vector<deferred_ptr<Node>>
            children {my_heap}, parents {my_heap};
        /*... data ...*/
    };
    vector<deferred_ptr<Node>> roots;
/*...*/
};
```



67

## Example: N-ary Graph

- ▶ **Q:** What if you want an N-ary Graph (not just left and right children)?

- ▶ **(experimental) A:** `vector<T, deferred_allocator<T>>`

- ▶ Convenience alias: `deferred_vector<T>`

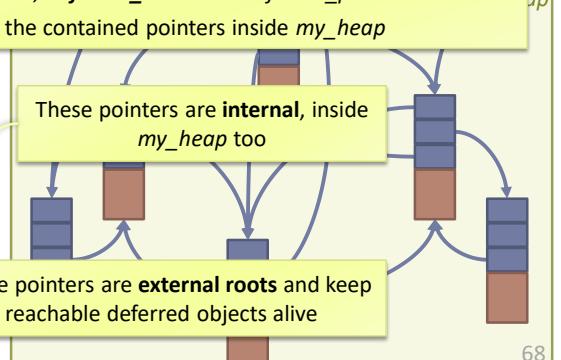
```
class Graph {
    static deferred_heap my_heap;
    struct Node {
        deferred_vector<deferred_ptr<Node>>
            children {my_heap}, parents {my_heap};
        /*... data ...*/
    };
    vector<deferred_ptr<Node>> roots;
/*...*/
};
```

`vector<deferred_ptr<Node>, deferred_allocator<deferred_ptr<Node>>>`

Used to keep all the contained pointers inside `my_heap`

These pointers are **internal**, inside `my_heap` too

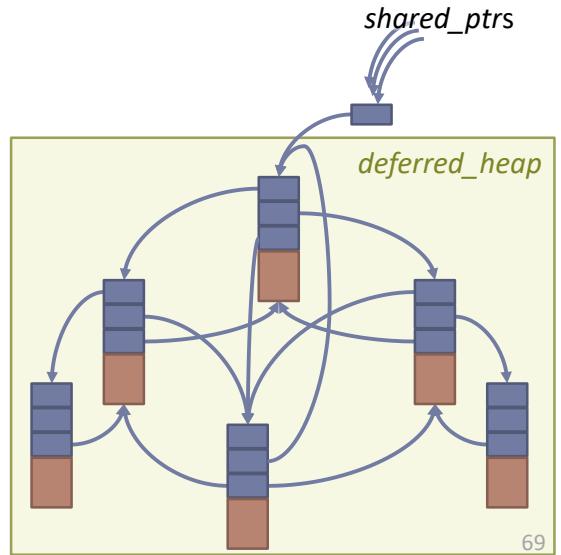
These pointers are **external roots** and keep reachable deferred objects alive



68

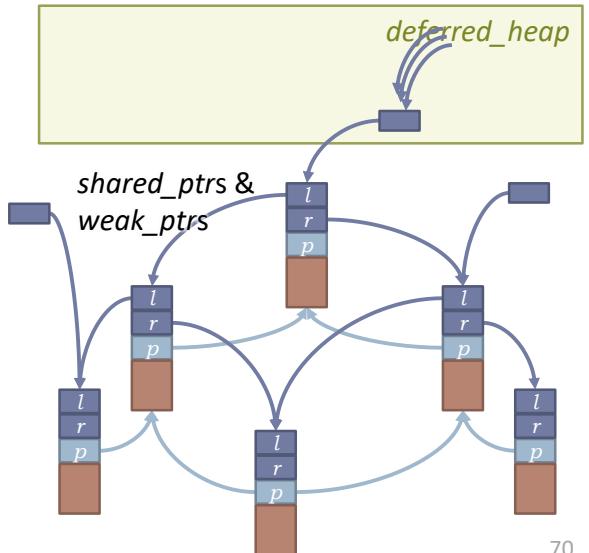
## Composition: RC lifetime → deferred lifetime

- Consider a graph of deferred objects rooted in an RC'd object.
  - The entire graph has **deterministic** RC lifetime, and gets **ordered** cleanup w.r.t. other such graphs.
  - The deferred objects nevertheless get to easily have **cycles** w.r.t. each other.
- Examples:
  - A *Graph* object and its internal nodes.
  - Region** lifetime = deallocate-at-once, but with real object destruction.



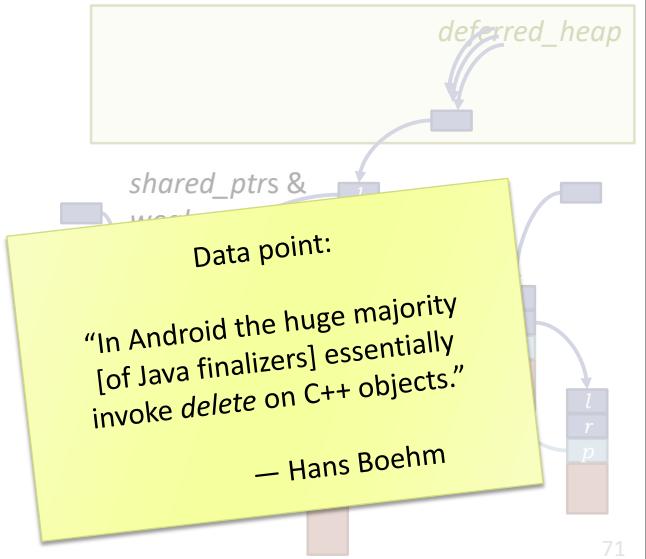
## Composition: Deferred lifetime → RC lifetime

- Consider a graph of RC'd objects rooted in a deferred object.
  - The entire graph has **lazy** lifetime, and can easily participate in **cycles** with other such graphs in the same *deferred\_heap*.
  - The RC objects nevertheless get **ordered** cleanup w.r.t. each other, and can call each other.
- Example:
  - A *Tree* object with lazy lifetime, but ordered internal node destruction within the tree.



## Composition: Deferred lifetime → RC lifetime

- ▶ Consider a graph of RC'd objects rooted in a deferred object.
- ▶ The entire graph has **lazy** lifetime, and can easily participate in **cycles** with other such graphs in the same **deferred\_heap**.
- ▶ The RC objects nevertheless get **ordered** cleanup w.r.t. each other, and can call each other.
- ▶ Example:
  - ▶ A *Tree* object with lazy lifetime, but ordered internal node destruction within the tree.



71

## `std::*_ptr` limitations

- ▶ Ownership cycles ⇒ **leak** objects and memory.
  - ▶ `shared_ptr` uses reference counting.
  - ▶ Plain reference counting leaks cycles.
  - ▶ Can't always be naturally broken using a `weak_ptr`.
- ▶ Unbounded pointer copy cost ⇒ exceed **real-time** execution deadlines.
  - ▶ `shared_ptr` calls destructors promptly. (Usually a great thing!)
  - ▶ But objects can own other objects transitively.
  - ▶ Plain assignment to a `shared_ptr` can have arbitrarily high cost. (Similar issue with EH.)
- ▶ Unbounded destructor depth ⇒ exceed **stack limits** in constrained environments.
  - ▶ Scoped/`unique_ptr`/`shared_ptr` lifetime calls destructors recursively (nested).
  - ▶ As above, ownership trees can be deep.
  - ▶ Releasing a tree of objects requires having sufficient stack space.

72

## *std::\*ptr limitations*

- ▶ Ownership cycles  $\Rightarrow$  leak objects and memory.
  - ▶ *shared\_ptr* uses reference counting.
- Key word: **deferred destruction**
  - ▶ Preventing leaks cycles.
  - ▶ Usually broken using a *weak\_ptr*.
- ▶ Unbounded pointer copy cost  $\Rightarrow$  exceed **real-time** execution demands
  - ▶ *shared\_ptr* calls destructors promptly. (Usually a great thing!)
  - ▶ But objects can own other objects transitively.
  - ▶ Plain assignment to a *shared\_ptr* can have arbitrarily high cost. (Similar issue with EH.)
- ▶ Unbounded destructor depth  $\Rightarrow$  exceed **stack limits** in constrained environments
  - ▶ *Shared/unique\_ptr/shared\_ptr* lifetime calls destructors recursively up trees can be deep.
  - ▶ Releasing a tree of objects requires having sufficient stack space.

Assigning a  
*deferred\_ptr* doesn't  
run destructors  
 $\Rightarrow$  predictable cost

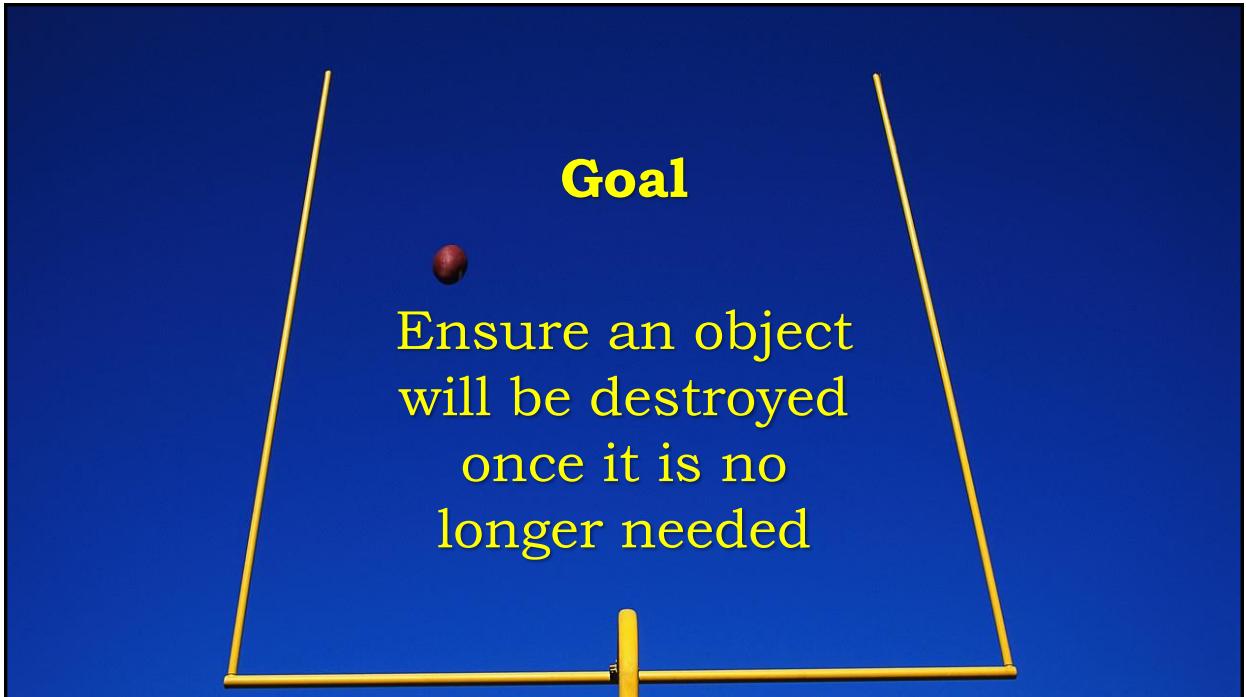
*deferred\_heap*  
::collect() runs  
destructors iteratively  
 $\Rightarrow$  less stack depth

73

## Note: *deferred\_ptr* has limitations too!

- ▶ **More total cost** (and/but moved around in different places).
  - ▶ It moves work to different places, and does more total work/overhead than *unique\_ptr* or *shared\_ptr*.
  - ▶ Prefer *unique\_ptr* or *shared\_ptr* in that order where possible, as usual. You should be using them much more frequently than *deferred\_ptr*.
- ▶ **A demo** of a technique for deferred and unordered destruction.
  - ▶ Experimental code to generate ideas you can draw from.
  - ▶ Not production quality.

74



## Leak freedom: In order of preference

Strategy	Natural examples	Cost	Rough frequency
<b>1. Prefer scoped lifetime by default</b> (locals, members)	Local and member objects – directly owned	Zero: Tied directly to another lifetime	O(80%) of objects
<b>2. Else prefer make_unique &amp; unique_ptr or a container</b> , if the object must have its own lifetime (i.e., heap) and ownership can be unique w/o owning cycles	Implementations of trees, lists	Same as new/delete & malloc/free <b>Automates</b> simple heap object use in a library	O(20%) of objects
<b>3. Else prefer make_shared &amp; shared_ptr</b> , if the object must have its own lifetime (i.e., heap) and shared ownership w/o owning cycles	Node-based DAGs, incl. trees that share out references	Same as manual reference counting (RC) <b>Automates</b> shared object use in a library	O(1%) of objects
<b>4. Else defer destruction</b> and have logic to trace unreachable objects (maybe eventually automate in a library?)	Graphs, real-time code, limited stacks	Deferred destructors and tracing logic	

## Leak Freedom in C++... *By Default*

**Questions?**



## FAQs

- ▶ Q: Is this garbage collection?
  - ▶ Of course, as is reference counting (e.g., `shared_ptr`).
- ▶ Q: I meant, is this like tracing GC in other popular languages?
  - ▶ No. It collects objects, not garbage.
  - ▶ Also: It's granular and opt-in, not global and default.
- ▶ Q: Is this related/similar to the Boehm collector?
  - ▶ No. It runs destructors, and tracing is granular (scoped to an individual heap) and accurate (not conservative).
- ▶ Q: Is this like a region?
  - ▶ It's a superset.
  - ▶  $\text{deferred\_heap} = \text{region} + \text{collect}(\text{()}_\text{opt}) + \text{destructors} + \text{nulling roots}$ .

81

## *deferred\_heap* requirements & constraints

### Do this... (appropriate for C++)

- ▶ **Opt-in:** Not the primary allocator, available to use just when you need it
- ▶ **Zero-overhead == “don’t pay for what you don’t use”:** Cost  $\propto$  GC allocations only
- ▶ **Deterministic:** Collection scope + timing
- ▶ **Real destructors:** No separate “finalizer” concept needed
- ▶ **Solid resource cleanup:** Enforce “don’t access other nontrivially destructible objects during cleanup” (can’t write it), prevent “resurrection” (can’t write it)

### ... not that (ok for other languages)

- ▶ Forced-in: Default/only allocator, must fight the language to get out
- ▶ Global cost (incl. “conservative” GC): Incur overhead on code/data that doesn’t use it
- ▶ Nondeterministic: Hard to control pauses
- ▶ Distinct-but-related finalizer: Leads to brittle “Dispose(bool)” patterns due to overlap
- ▶ Complex resource cleanup: Unenforced guidance to “don’t access other finalizable objects” guidance, allows re-reachability “resurrection”

82